

Lecture 7:

Optimizations

Modern Parallel Computing
John Owens
EEC 289Q, UC Davis, Winter 2018

based on lecture by Kerry Seitz

Credits

- **Thanks to John Stratton and Anjul Patney for providing slides**

**J. A. Stratton, C. Rodrigues, I.-J. Sung, L.-W. Chang, N. Anssari, G. Liu, W.-M. W. Hwu, and N. Obeid,
“Algorithm and Data Optimization Techniques for Scaling to Massively Threaded Systems,” Computer,
vol. 45, no. 8, pp. 26–32, 2012.**

Announcements


- **HW2: Suggest you use `cudaMallocManaged` and linearize your input/output matrices**
- **Guest speakers next 3 Thursdays**
 - **Please come on time**
 - **Office hours will be perturbed, feel free to ask for alternate arrangements, will keep you posted**

Basic Efficiency Rules

- **Develop algorithms with a data parallel mindset**
- **Minimize divergence of execution within blocks**
- **Maximize locality of global memory accesses**
 - **“Coalescing”**
- **Exploit per-block shared memory as scratchpad**
- **Expose enough parallelism**

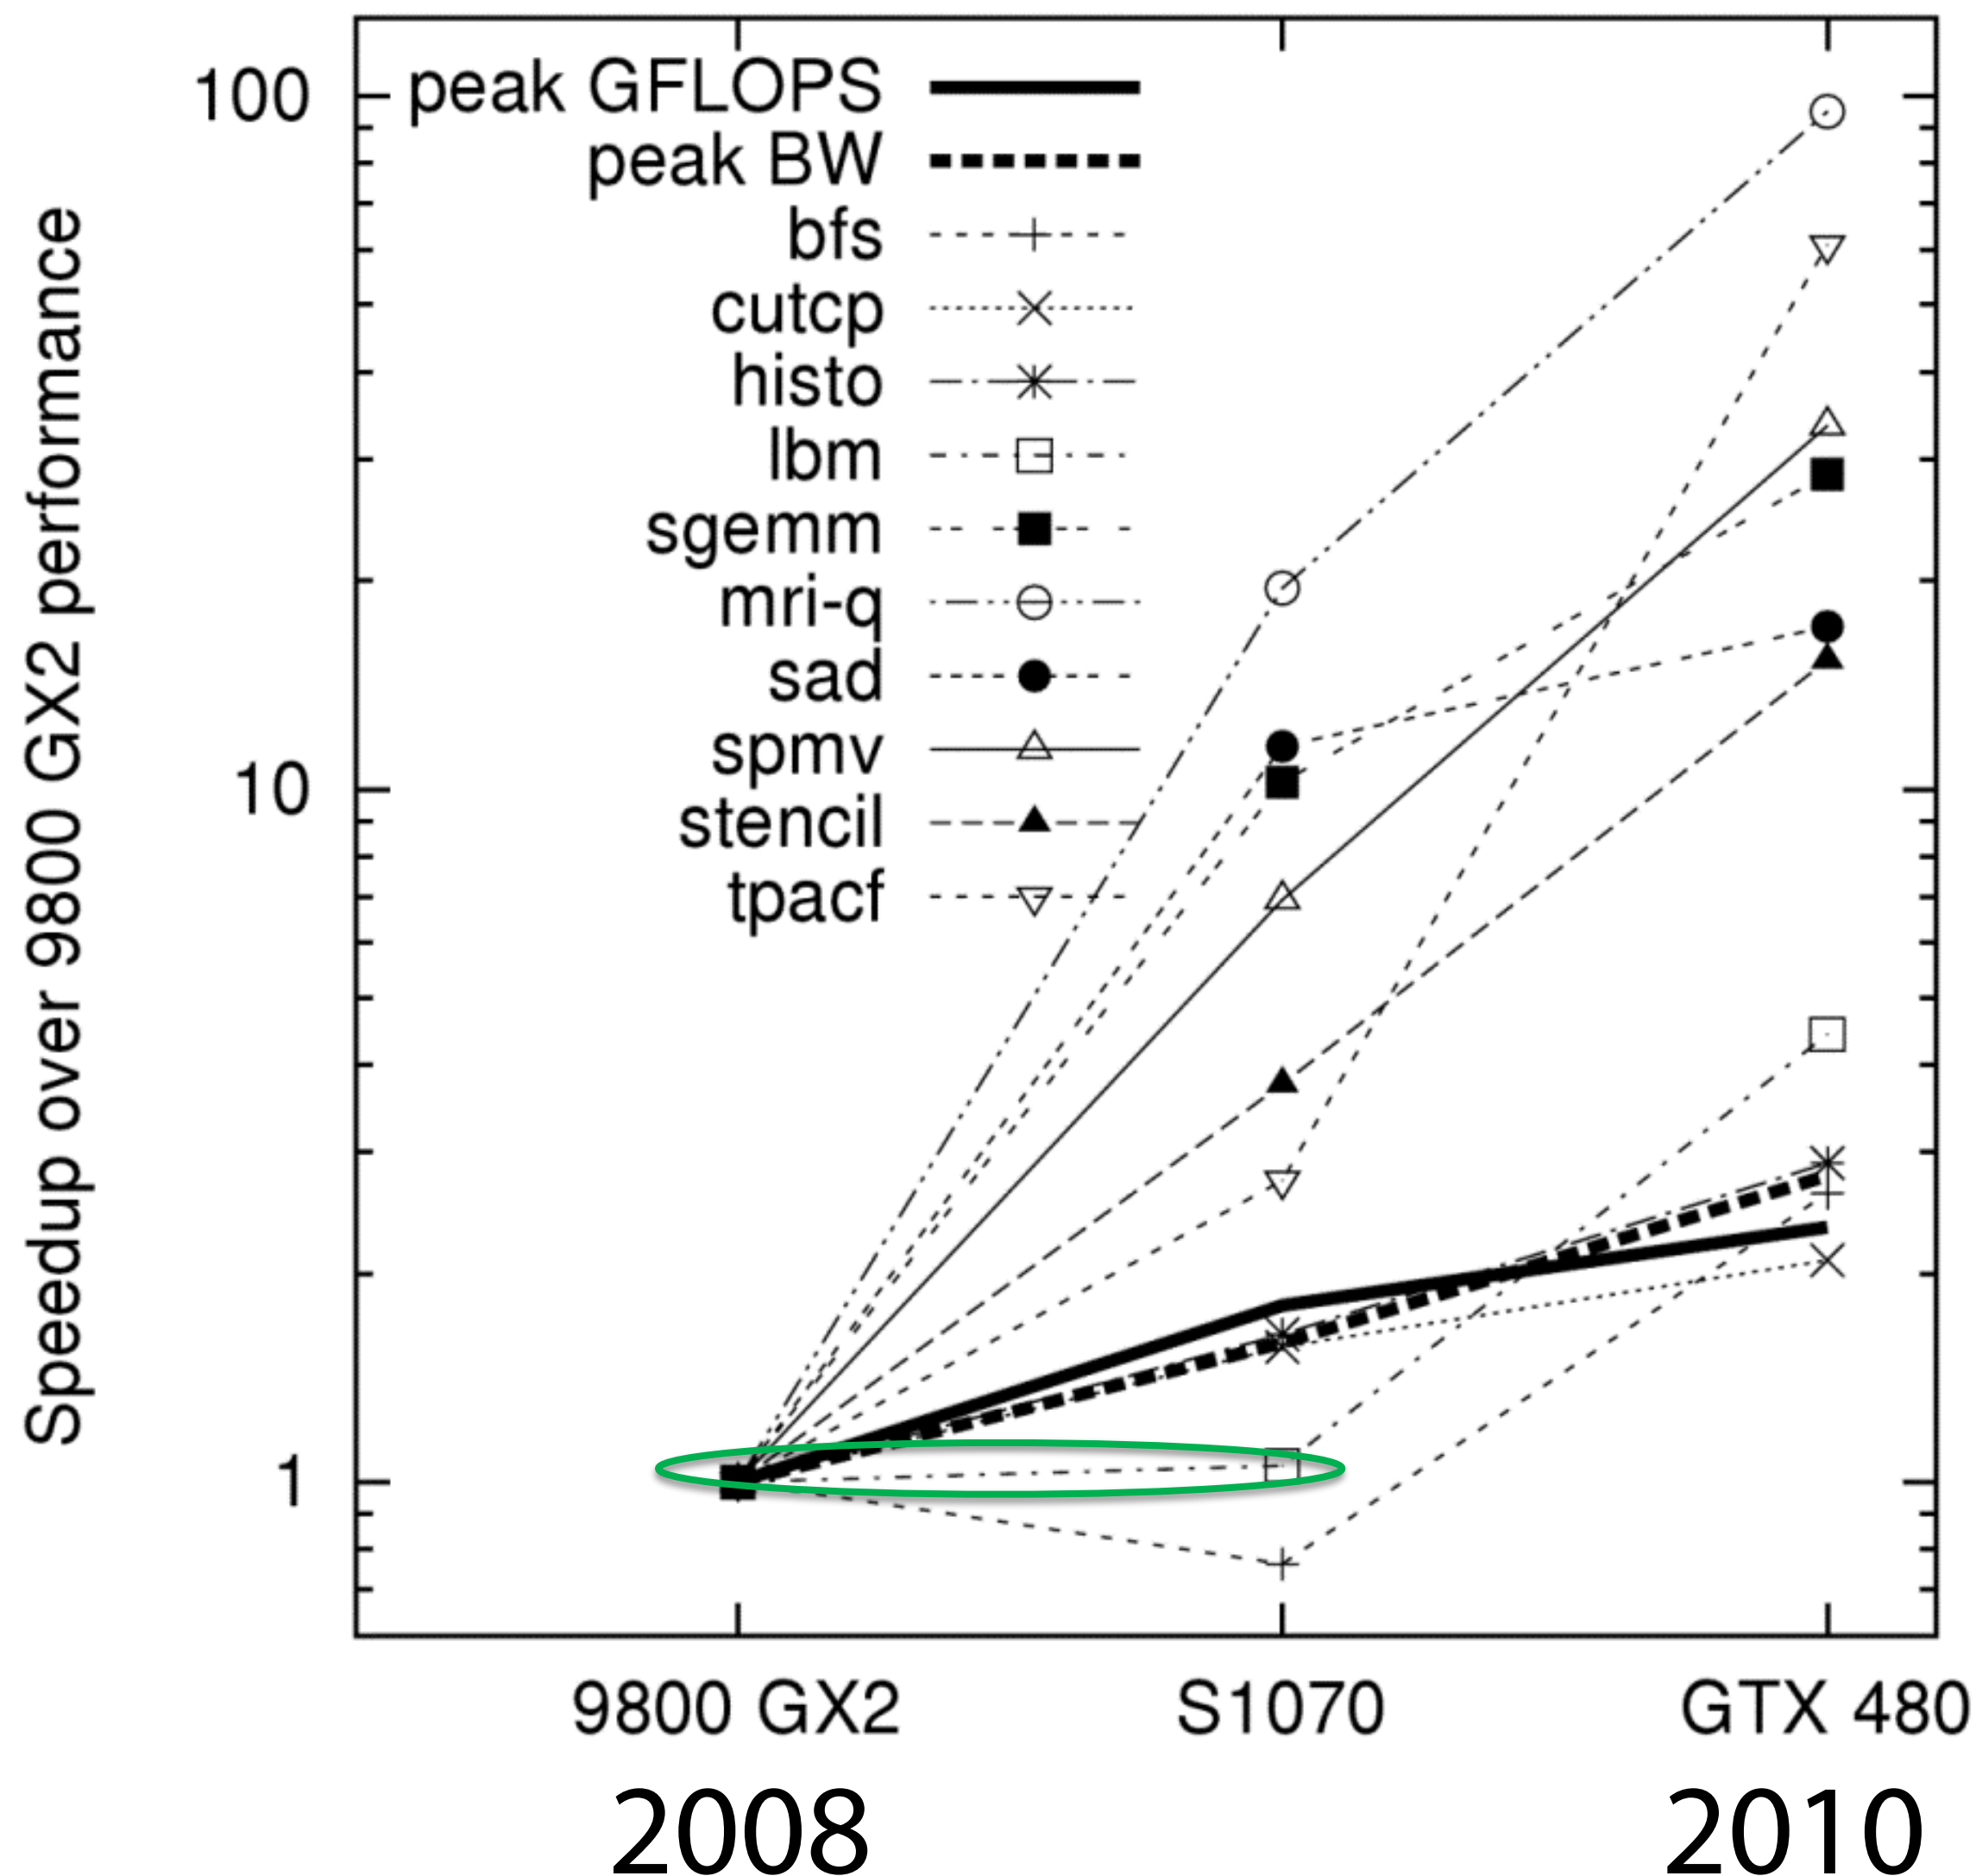
**How much faster do
applications really get each
hardware generation?**

Unoptimized Code Has Improved Drastically




- Orders of magnitude speedup in many cases
- Hardware does not solve all problems

Unoptimized Code Has Improved Drastically



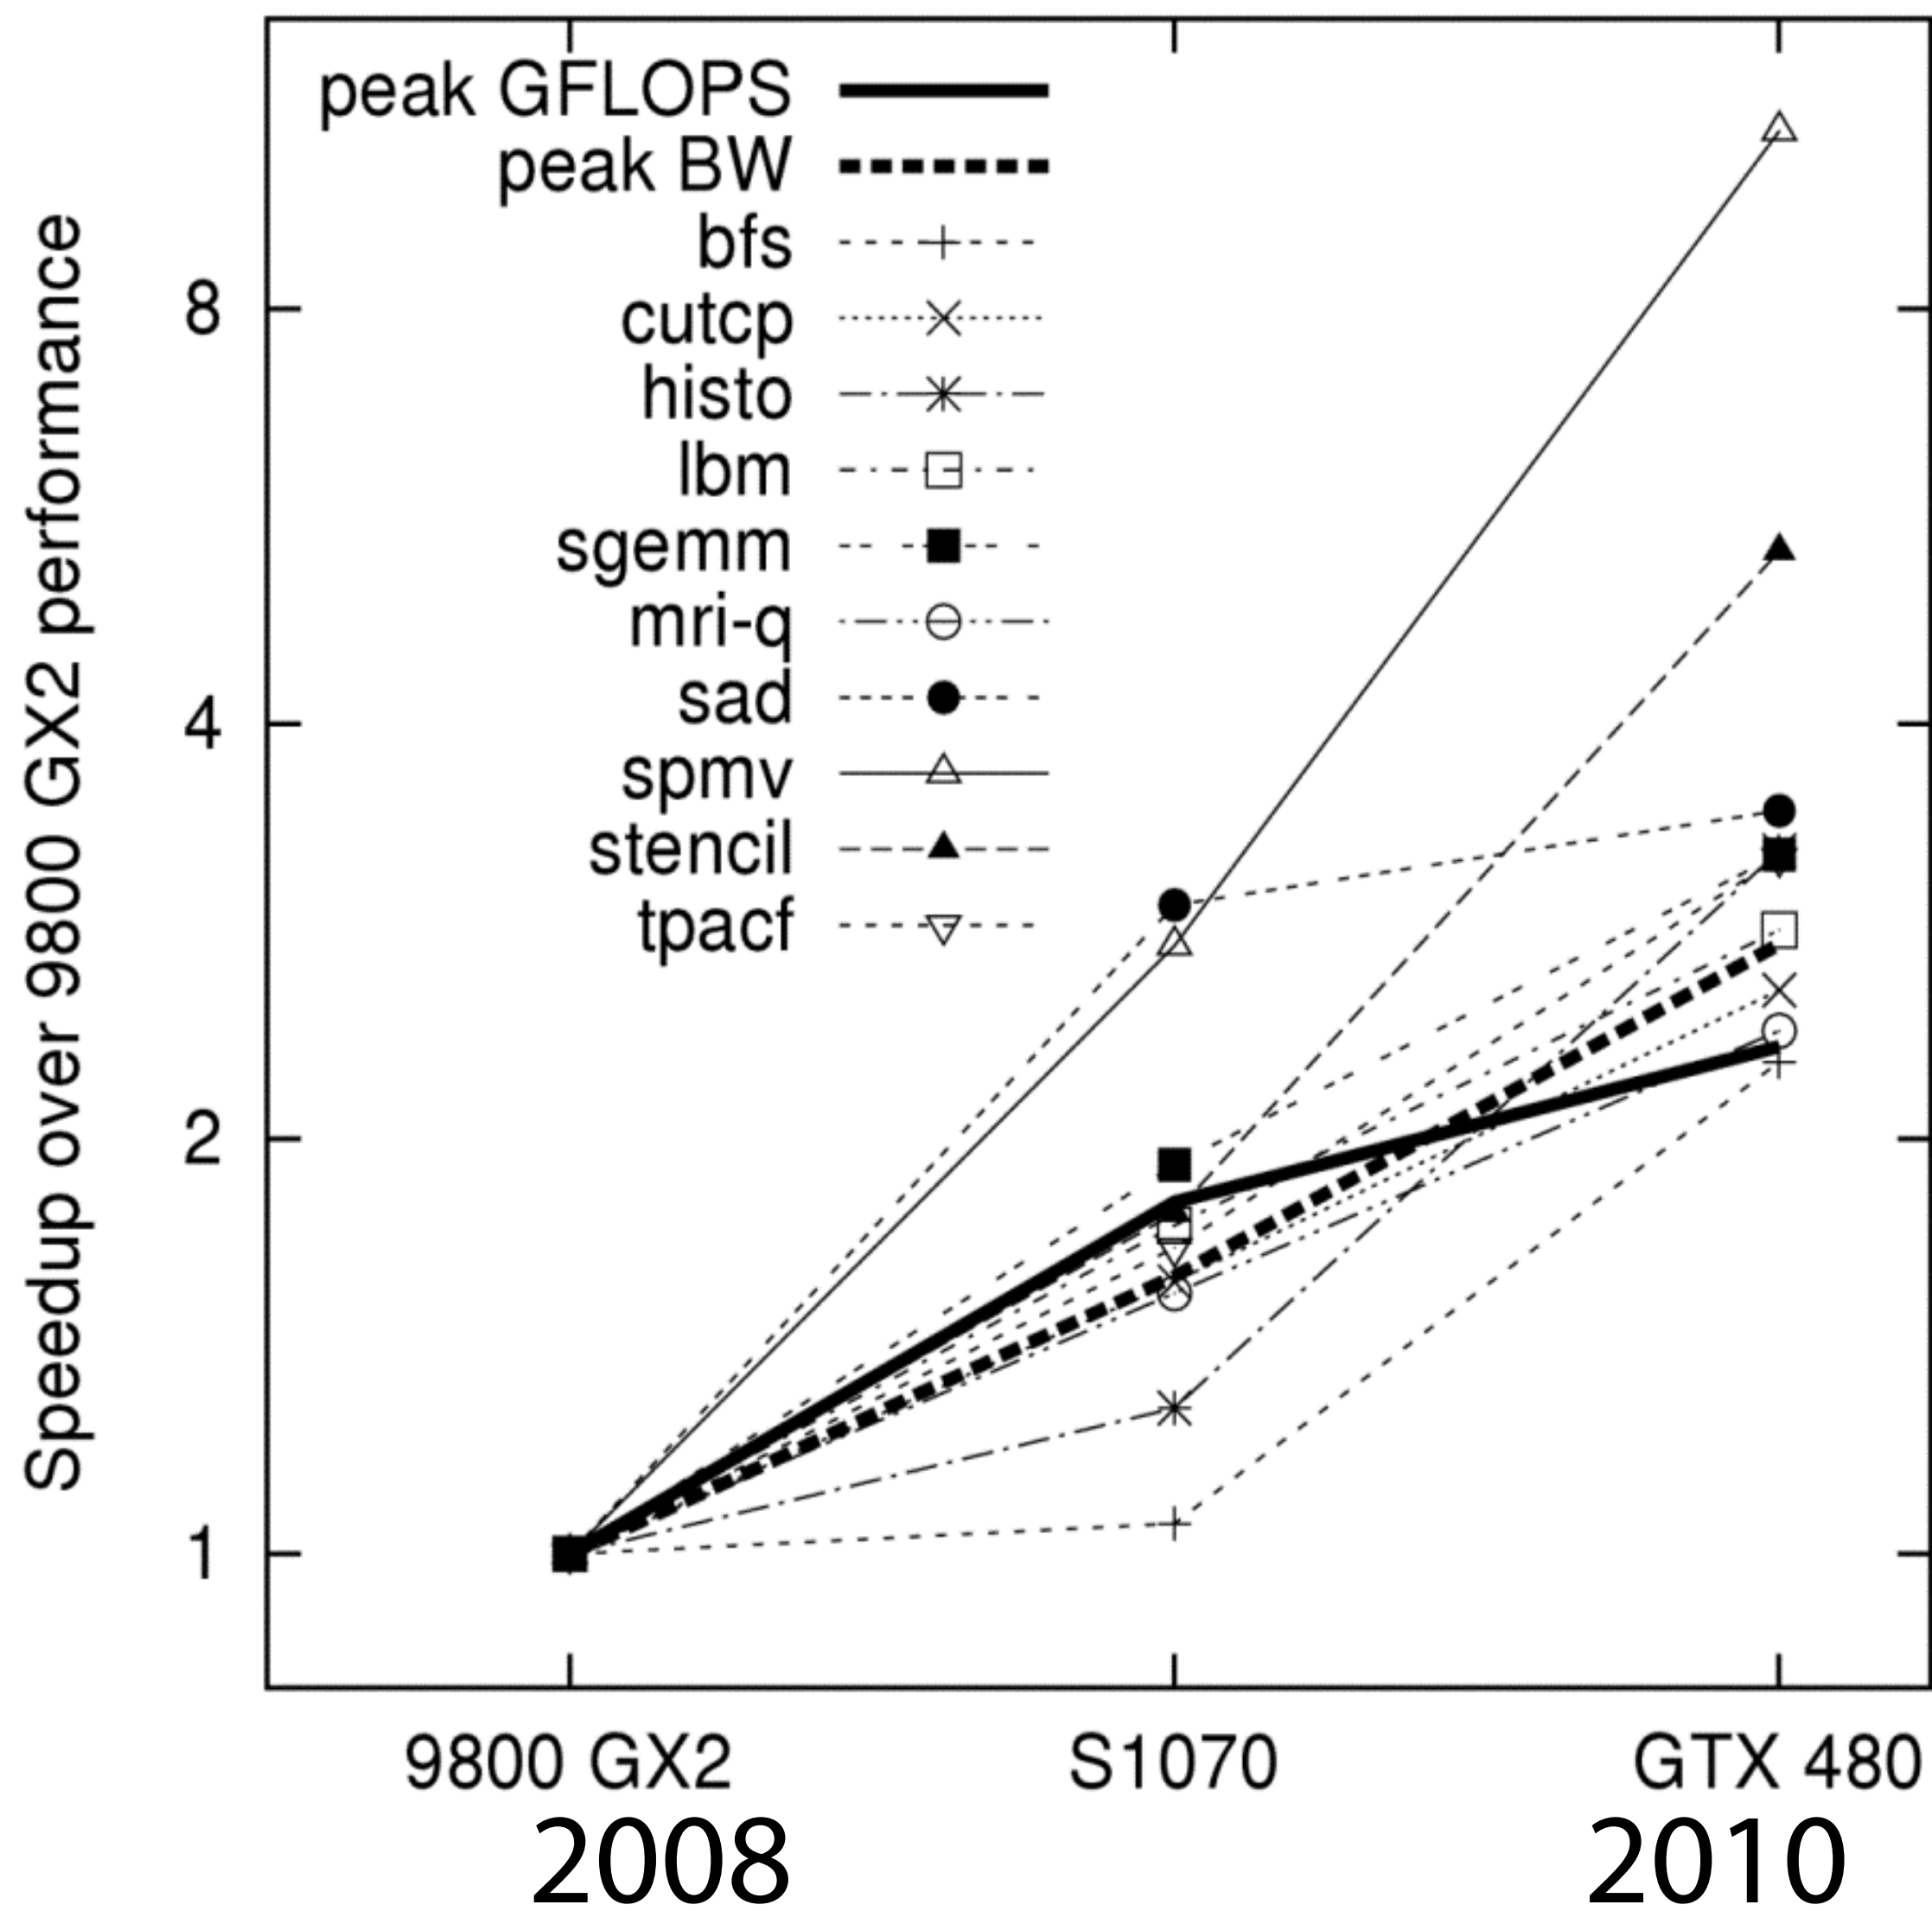
- Orders of magnitude speedup in many cases
- Hardware does not solve all problems
 - Coalescing (**lbm**)

Unoptimized Code Has Improved Drastically



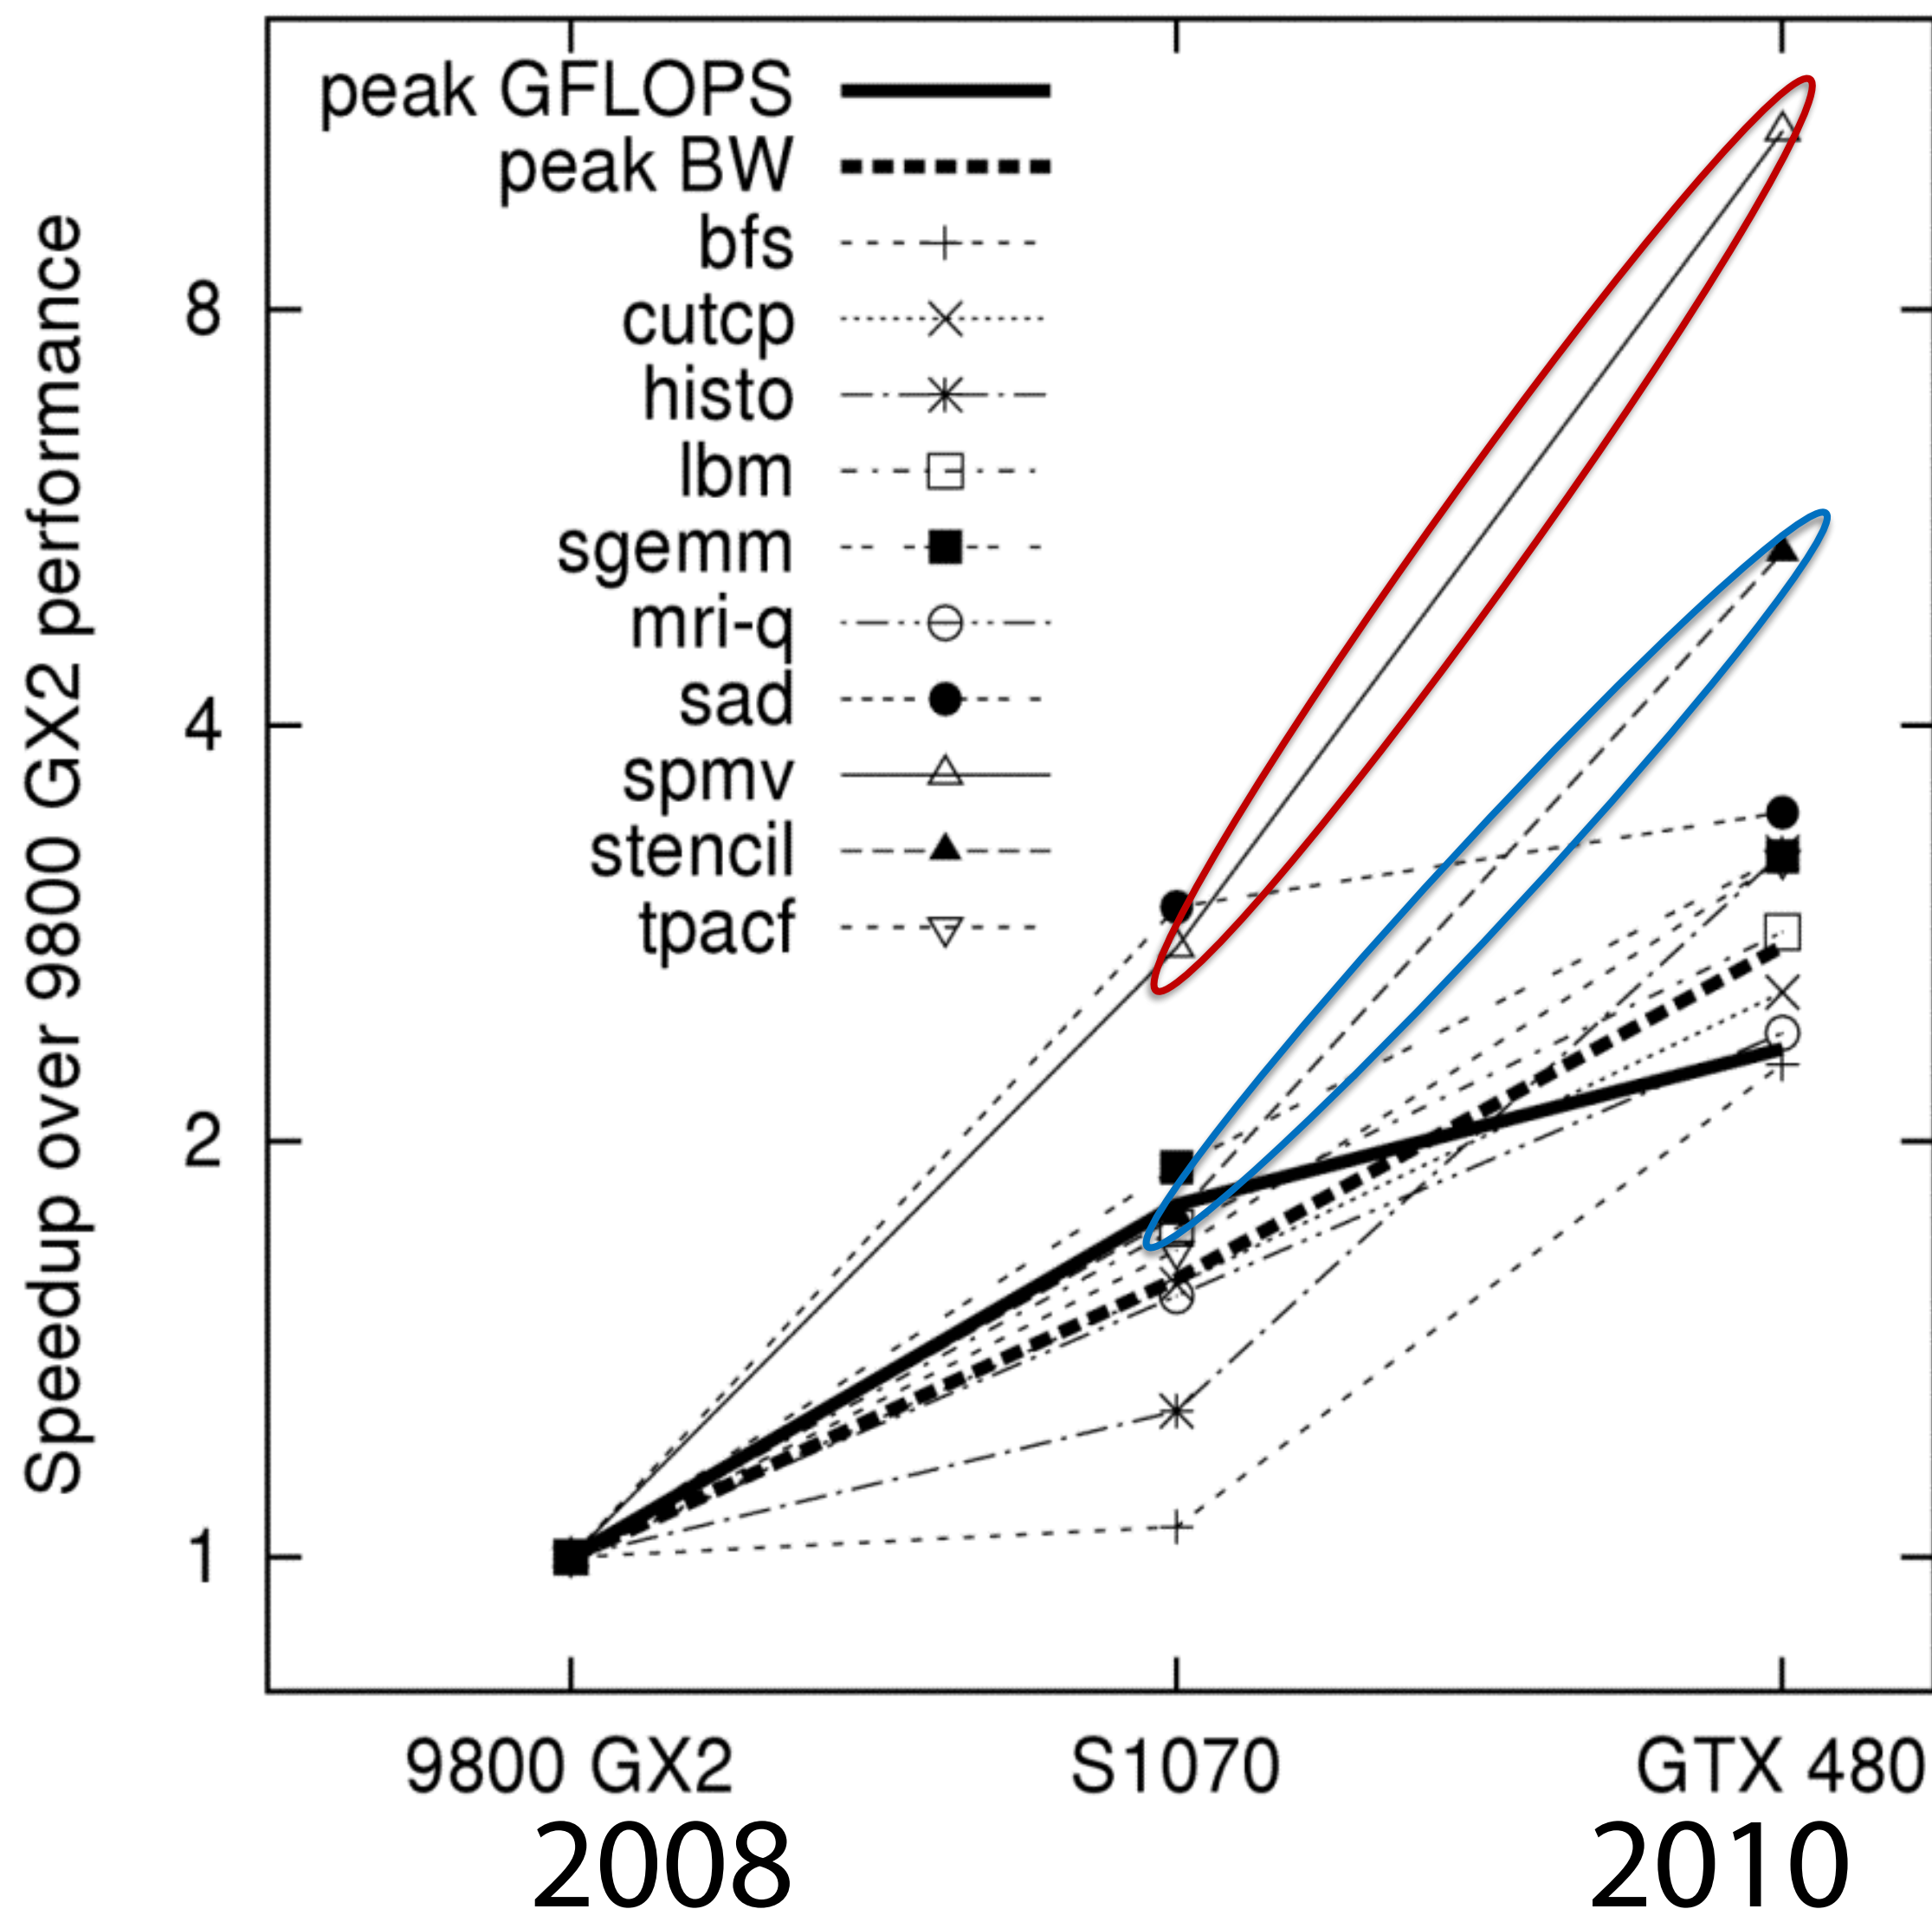
- Orders of magnitude speedup in many cases
- Hardware does not solve all problems
 - Coalescing (**lbm**)
 - Highly contentious atomics (**bfs**)

Optimized Code Is Improving Faster than “Peak Performance”




Optimized Code Is Improving Faster than “Peak Performance”

- Caches capture locality scratchpad can't efficiently (**spmv**, **stencil**)

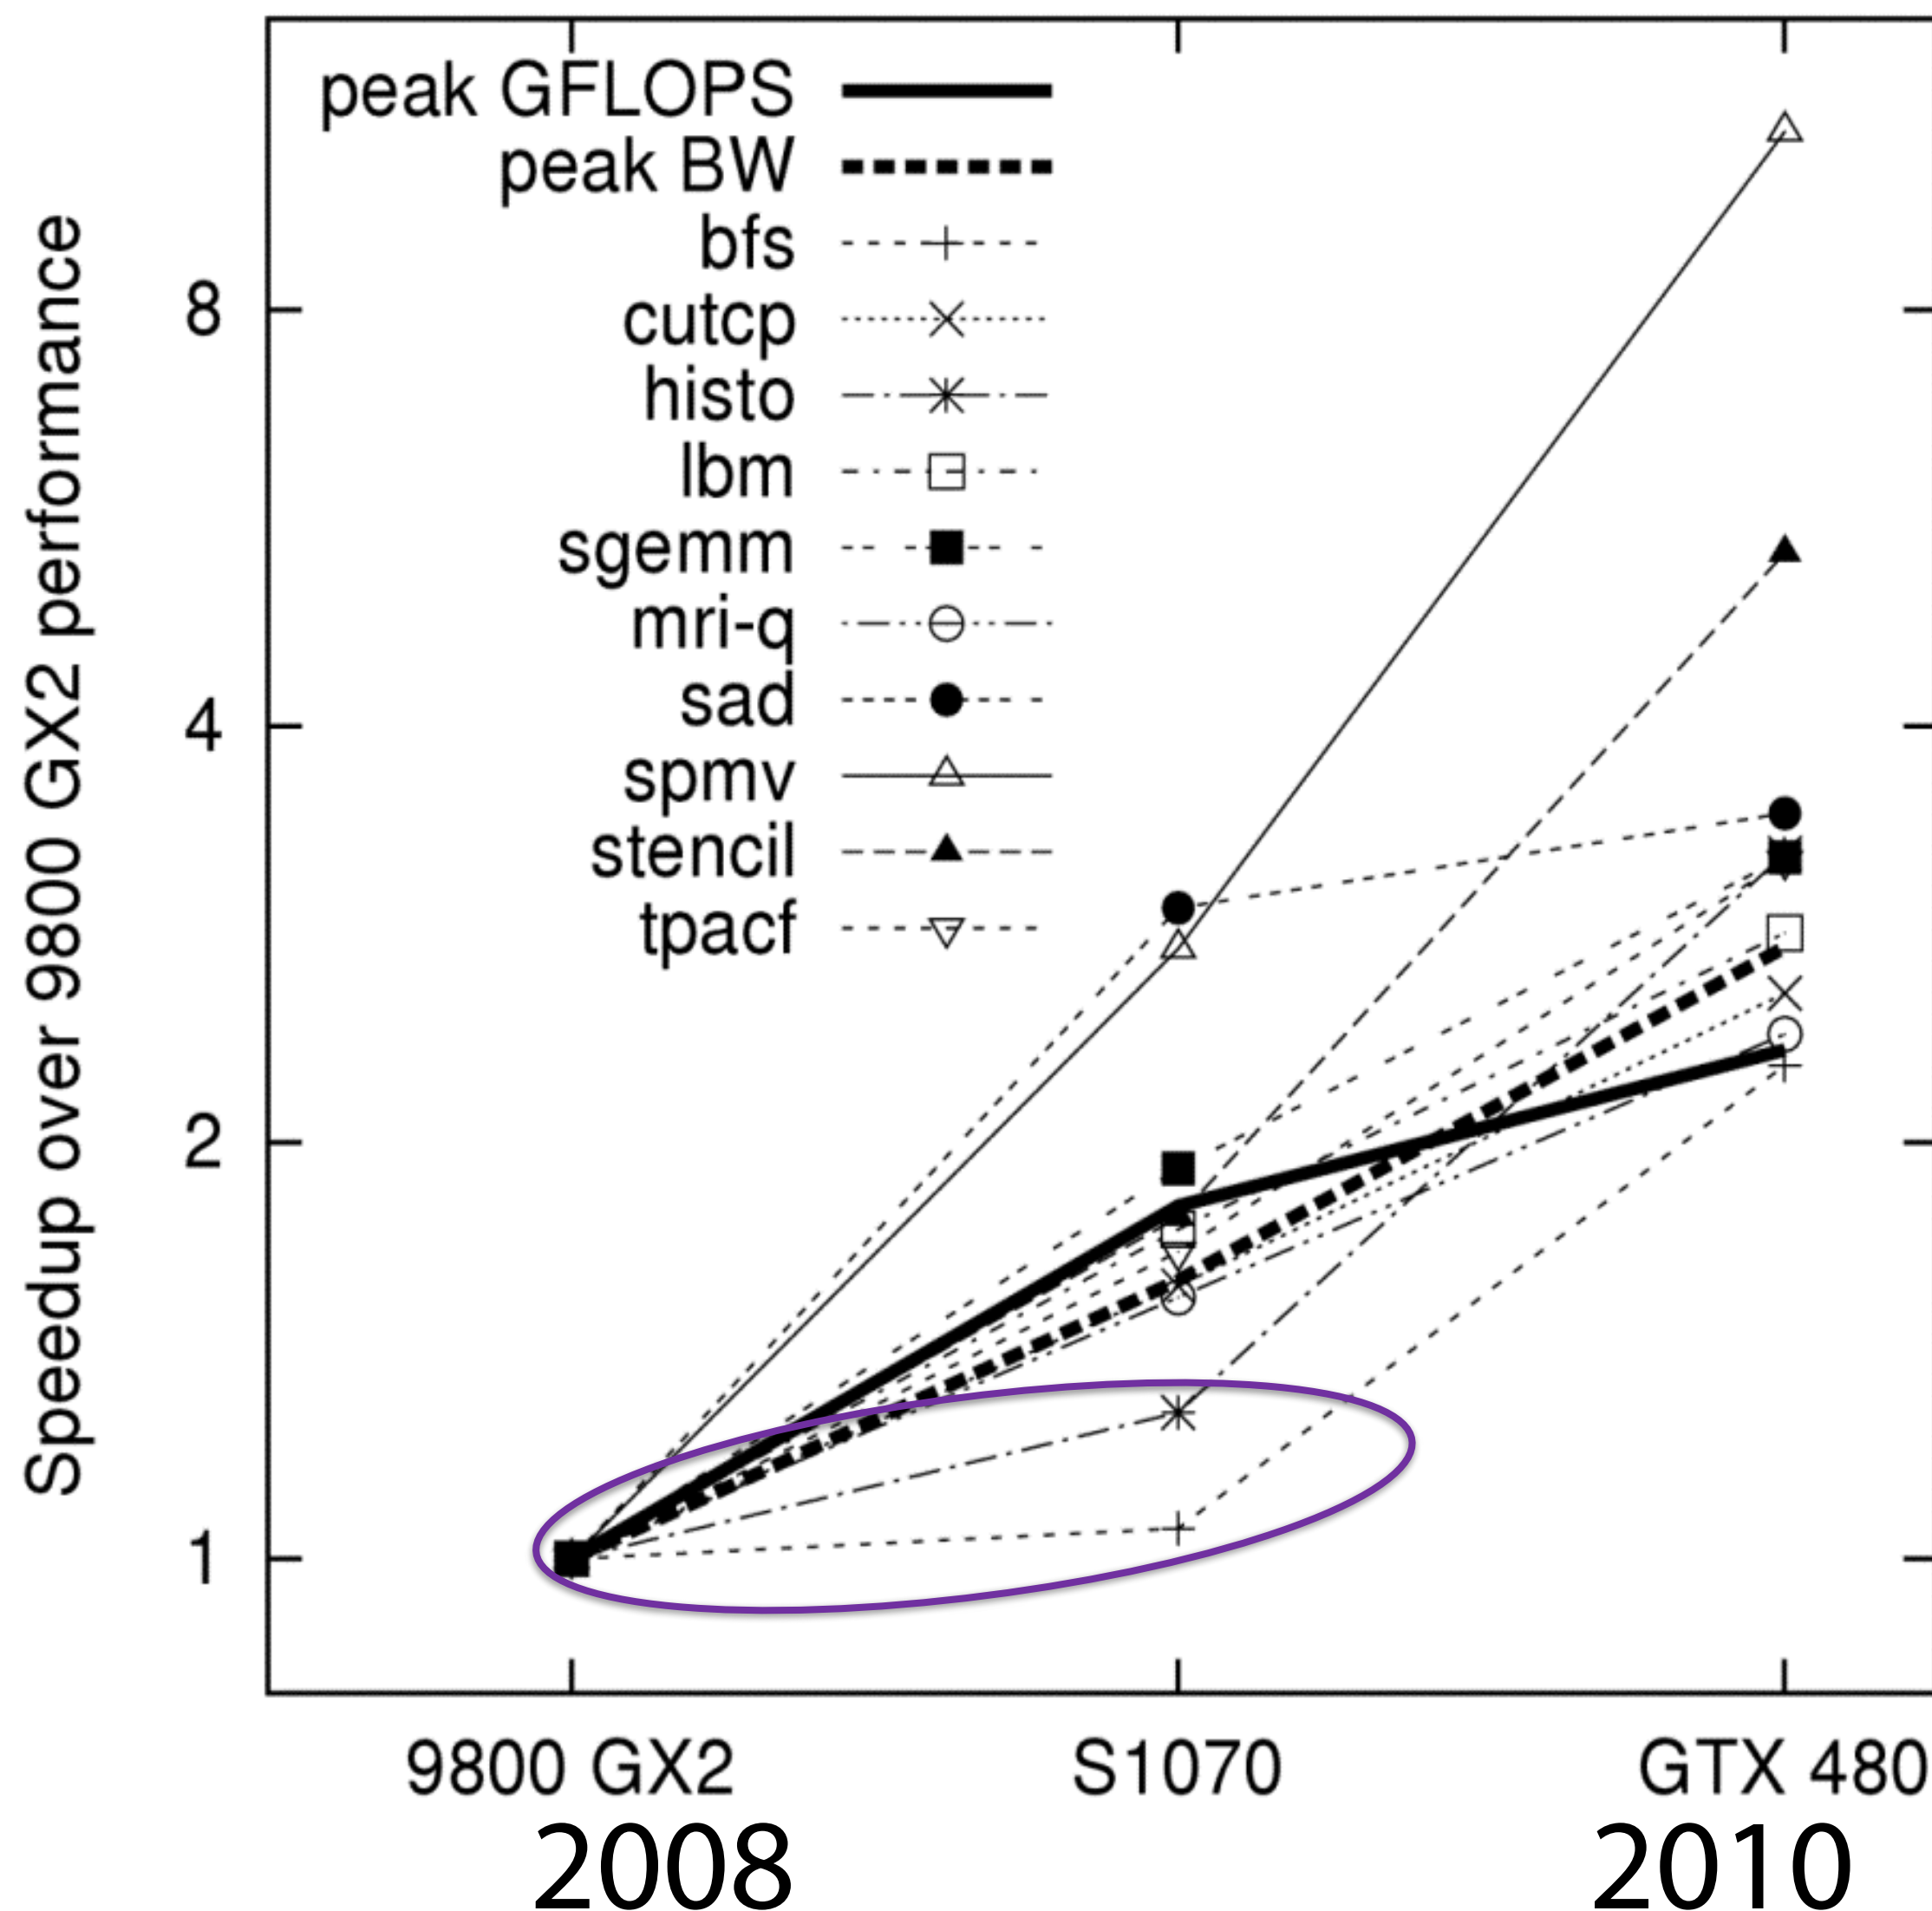


Optimized Code Is Improving Faster than “Peak Performance”



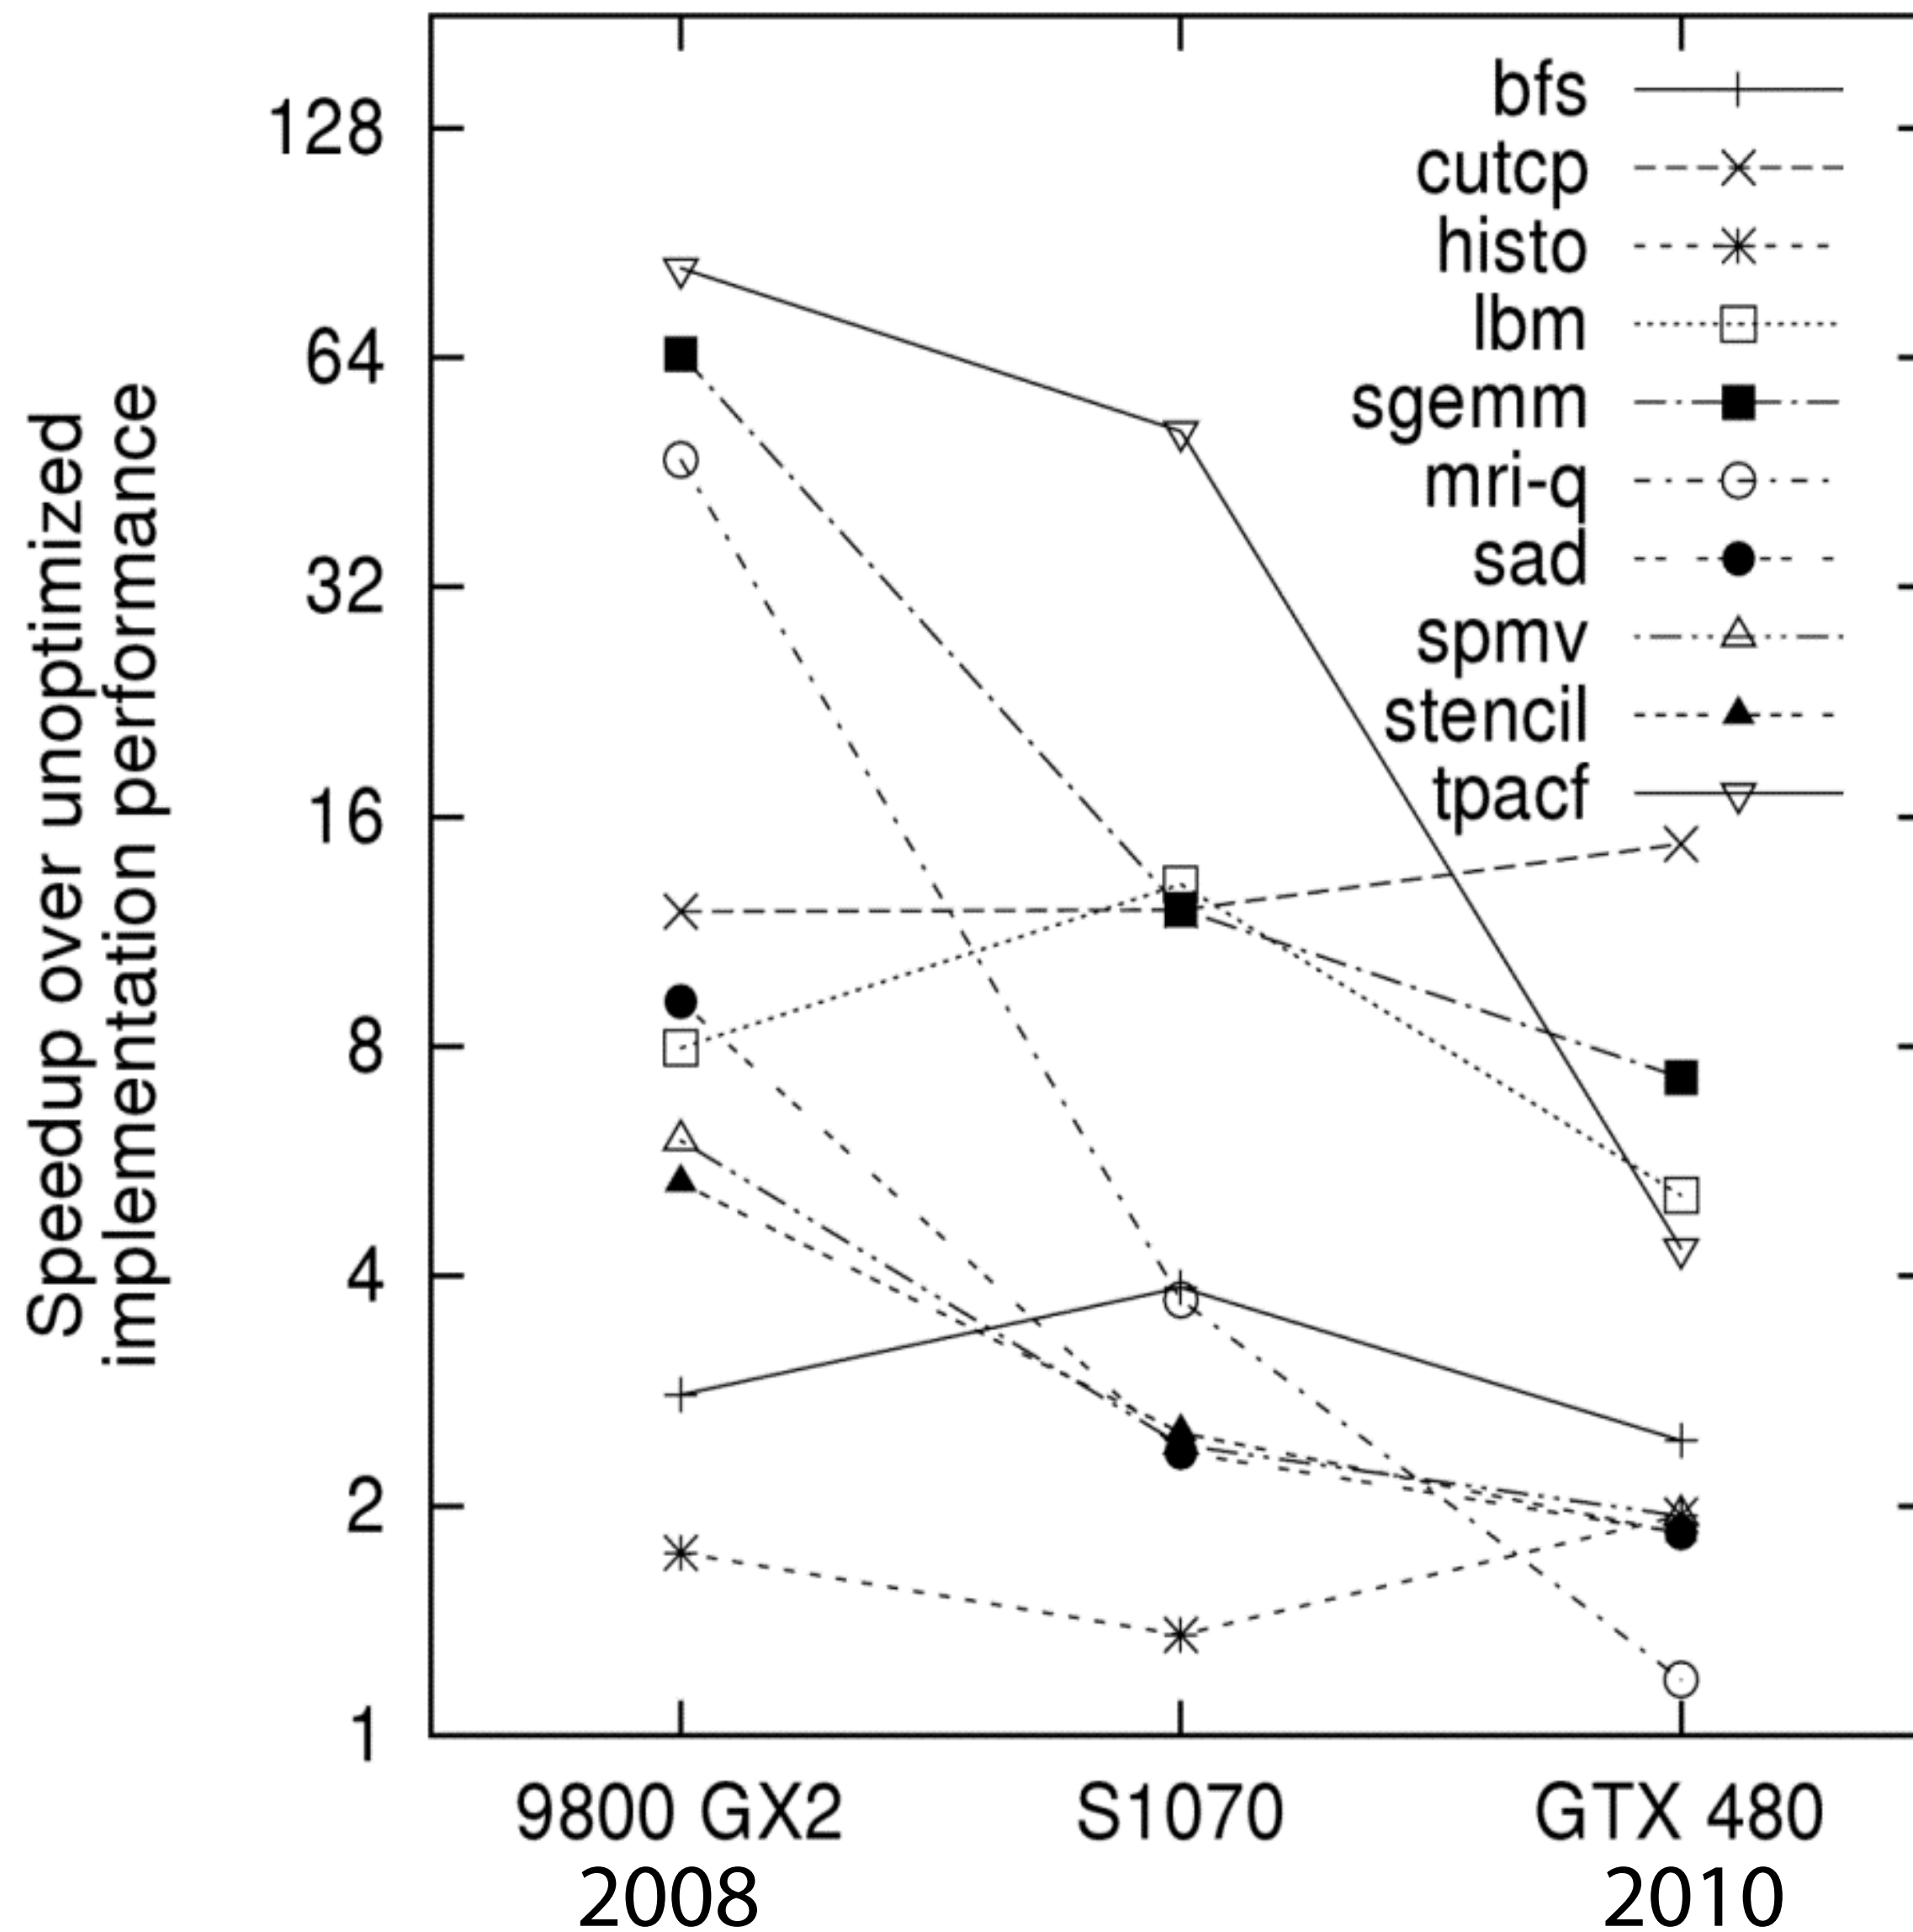
- Caches capture locality scratchpad can't efficiently (**spmv**, **stencil**)
- Increased local storage capacity enables extra optimization (**sad**)

Optimized Code Is Improving Faster than “Peak Performance”




- Caches capture locality scratchpad can't efficiently (**spmv**, **stencil**)
- Increased local storage capacity enables extra optimization (**sad**)
- Some benchmarks need atomic throughput more than flops (**bfs**, **histo**)

Optimization Still Matters

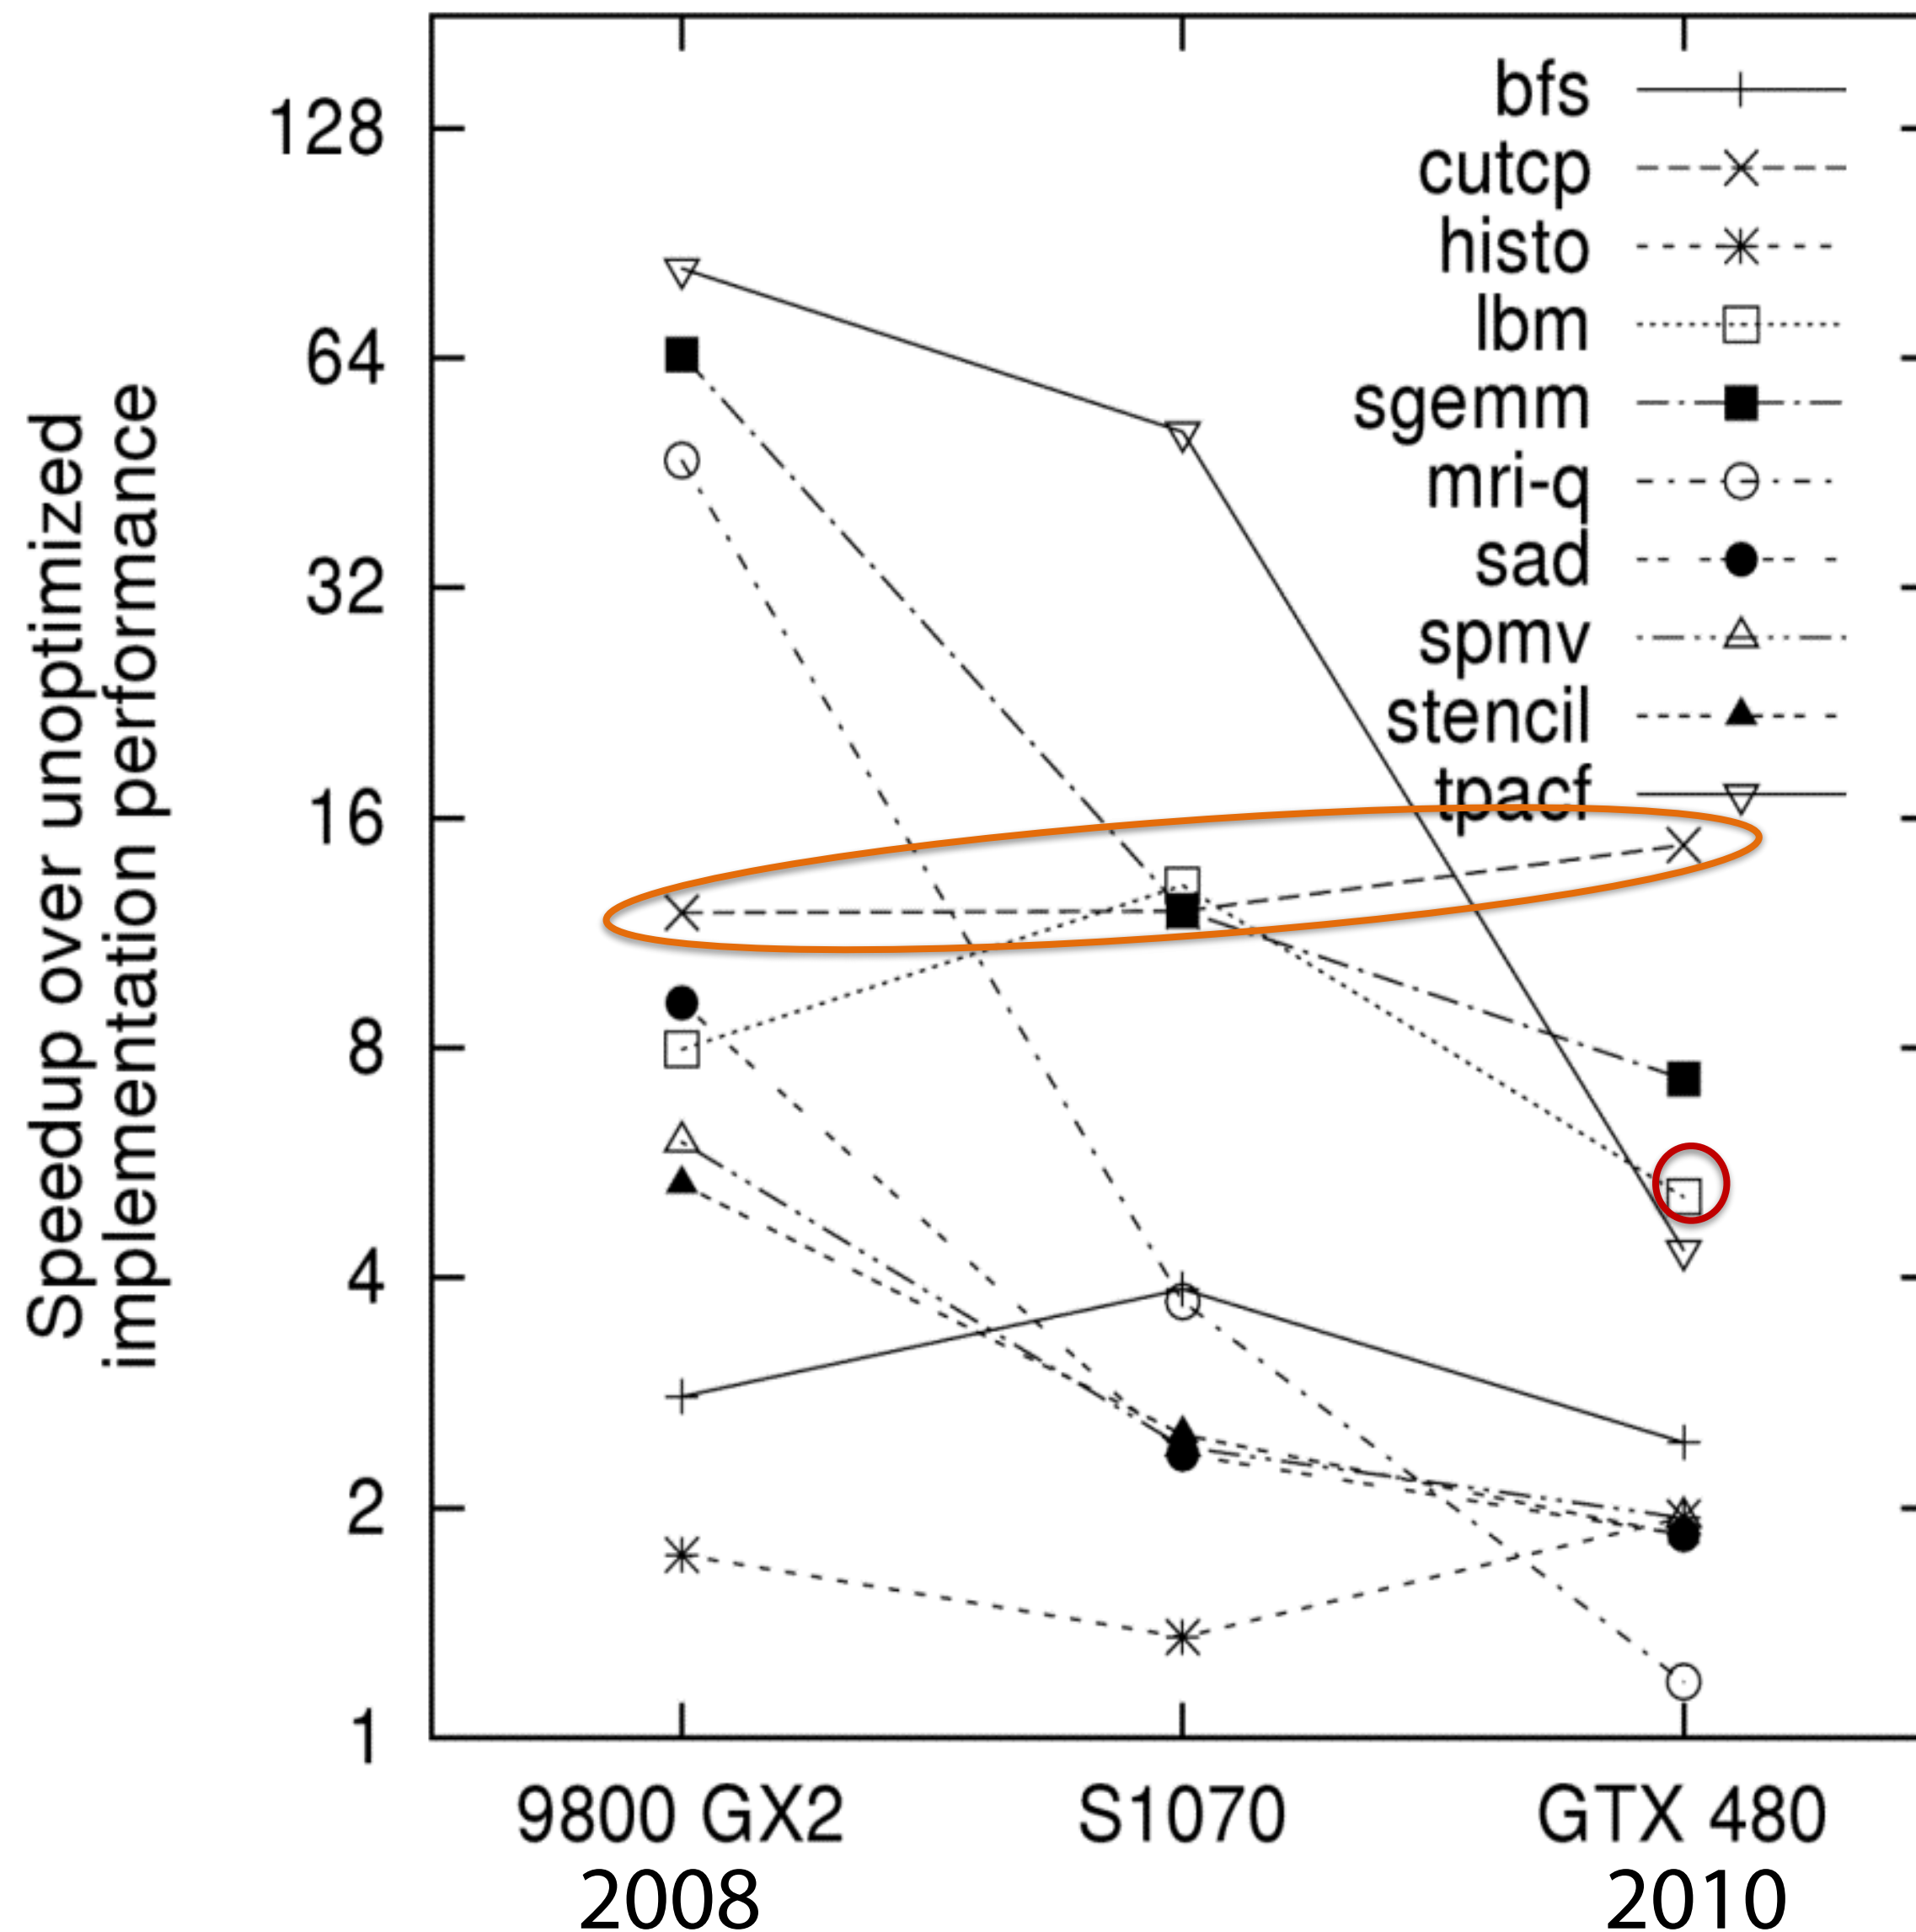


Optimization Still Matters




- Hardware never changes algorithmic complexity (**cutcp**)

Optimization Still Matters



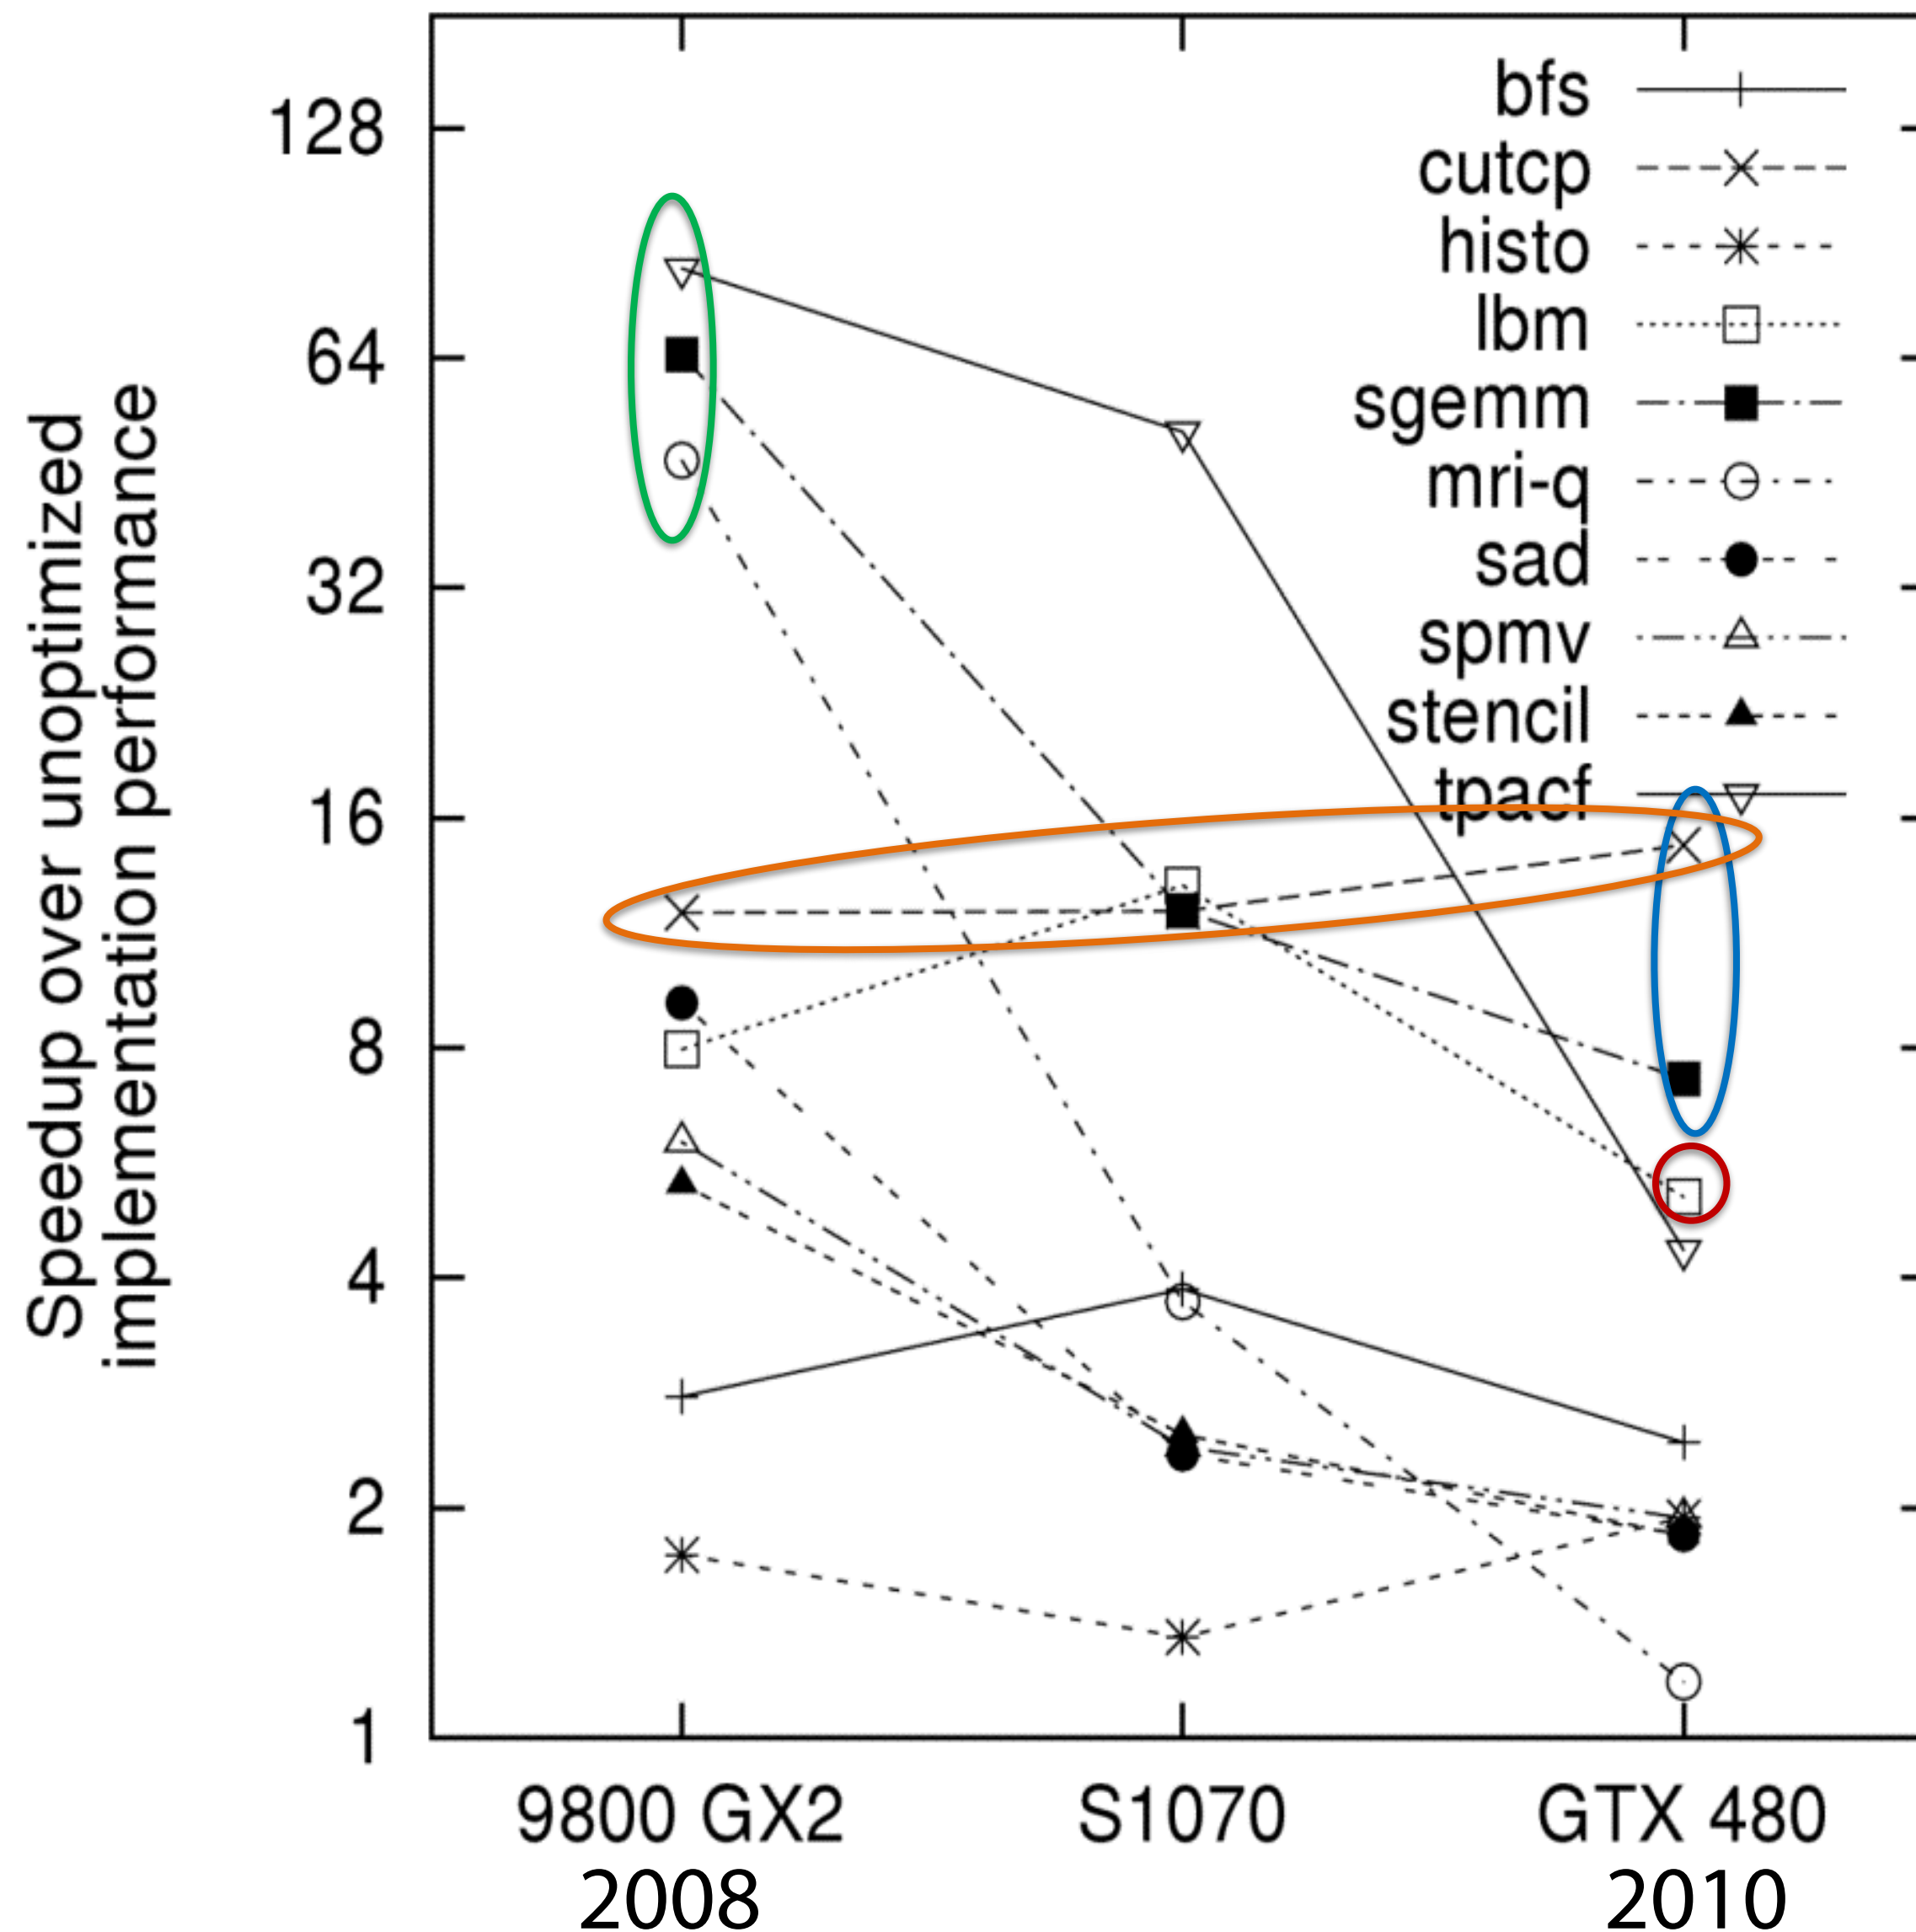
- Hardware never changes algorithmic complexity (**cutcp**)
- Caches do not solve layout problems for big data (**lbm**)

Optimization Still Matters



- Hardware never changes algorithmic complexity (**cutcp**)
- Caches do not solve layout problems for big data (**lbm**)
- Coarsening still makes a big difference (**cutcp**, **sgemm**)

Optimization Still Matters



- Hardware never changes algorithmic complexity (**cutcp**)
- Caches do not solve layout problems for big data (**lbm**)
- Coarsening still makes a big difference (**cutcp**, **sgemm**)
- Many artificial performance cliffs are gone (**sgemm**, **tpacf**, **mri-q**)

Speculations and Takeaways

Speculations and Takeaways

- **Optimizations still necessary today are unlikely to be magically solved by future hardware**
 - **Still necessary for highly parallel CPUs, after all**

Speculations and Takeaways

- **Optimizations still necessary today are unlikely to be magically solved by future hardware**
 - **Still necessary for highly parallel CPUs, after all**
- **Features matter just as much as FLOPS and GBytes/sec on lots of applications**
 - **Having a cache is critical, period**


Speculations and Takeaways

- **Optimizations still necessary today are unlikely to be magically solved by future hardware**
 - **Still necessary for highly parallel CPUs, after all**
- **Features matter just as much as FLOPS and GBytes/sec on lots of applications**
 - **Having a cache is critical, period**
- **Beware of unscalable implementation decisions**
 - **Global contention and synchronization will get worse**

Application Survey

- **Surveyed the GPU Computing Gems chapters**
- **Studied the Parboil benchmarks in detail**
- **Results:**
- **Nine (for now) major categories of optimization transformations**
 - **Performance impact of individual optimizations on certain Parboil benchmarks included in the paper**

Sparse matrix-vector multiplication (SpMV)



Sparse matrix-vector multiplication (SpMV)

$$\begin{pmatrix} \mathbf{A}_{0,0} & \mathbf{A}_{0,1} & \mathbf{A}_{0,2} & \mathbf{A}_{0,3} \\ \mathbf{A}_{1,0} & \mathbf{A}_{1,1} & \mathbf{A}_{1,2} & \mathbf{A}_{1,3} \\ \mathbf{A}_{2,0} & \mathbf{A}_{2,1} & \mathbf{A}_{2,2} & \mathbf{A}_{2,3} \\ \mathbf{A}_{3,0} & \mathbf{A}_{3,1} & \mathbf{A}_{3,2} & \mathbf{A}_{3,3} \end{pmatrix} \begin{pmatrix} \mathbf{x} \\ \mathbf{y} \\ \mathbf{z} \\ \mathbf{w} \end{pmatrix} = \begin{pmatrix} \mathbf{x}\mathbf{A}_{0,0} + \mathbf{y}\mathbf{A}_{0,1} + \mathbf{z}\mathbf{A}_{0,2} + \mathbf{w}\mathbf{A}_{0,3} \\ \mathbf{x}\mathbf{A}_{1,0} + \mathbf{y}\mathbf{A}_{1,1} + \mathbf{z}\mathbf{A}_{1,2} + \mathbf{w}\mathbf{A}_{1,3} \\ \mathbf{x}\mathbf{A}_{2,0} + \mathbf{y}\mathbf{A}_{2,1} + \mathbf{z}\mathbf{A}_{2,2} + \mathbf{w}\mathbf{A}_{2,3} \\ \mathbf{x}\mathbf{A}_{3,0} + \mathbf{y}\mathbf{A}_{3,1} + \mathbf{z}\mathbf{A}_{3,2} + \mathbf{w}\mathbf{A}_{3,3} \end{pmatrix}$$

Sparse matrix-vector multiplication (SpMV)

Row in input matrix \Rightarrow row in output vector

$$\begin{pmatrix} \mathbf{A}_{0,0} & \mathbf{A}_{0,1} & \mathbf{A}_{0,2} & \mathbf{A}_{0,3} \\ \mathbf{A}_{1,0} & \mathbf{A}_{1,1} & \mathbf{A}_{1,2} & \mathbf{A}_{1,3} \\ \mathbf{A}_{2,0} & \mathbf{A}_{2,1} & \mathbf{A}_{2,2} & \mathbf{A}_{2,3} \\ \mathbf{A}_{3,0} & \mathbf{A}_{3,1} & \mathbf{A}_{3,2} & \mathbf{A}_{3,3} \end{pmatrix} \begin{pmatrix} \mathbf{x} \\ \mathbf{y} \\ \mathbf{z} \\ \mathbf{w} \end{pmatrix} = \begin{pmatrix} \mathbf{x}\mathbf{A}_{0,0} + \mathbf{y}\mathbf{A}_{0,1} + \mathbf{z}\mathbf{A}_{0,2} + \mathbf{w}\mathbf{A}_{0,3} \\ \mathbf{x}\mathbf{A}_{1,0} + \mathbf{y}\mathbf{A}_{1,1} + \mathbf{z}\mathbf{A}_{1,2} + \mathbf{w}\mathbf{A}_{1,3} \\ \mathbf{x}\mathbf{A}_{2,0} + \mathbf{y}\mathbf{A}_{2,1} + \mathbf{z}\mathbf{A}_{2,2} + \mathbf{w}\mathbf{A}_{2,3} \\ \mathbf{x}\mathbf{A}_{3,0} + \mathbf{y}\mathbf{A}_{3,1} + \mathbf{z}\mathbf{A}_{3,2} + \mathbf{w}\mathbf{A}_{3,3} \end{pmatrix}$$

Sparse matrix-vector multiplication (SpMV)

Row in input matrix \Rightarrow row in output vector

$$\begin{pmatrix} \mathbf{A}_{0,0} & \mathbf{A}_{0,1} & \mathbf{A}_{0,2} & \mathbf{A}_{0,3} \\ \mathbf{A}_{1,0} & \mathbf{A}_{1,1} & \mathbf{A}_{1,2} & \mathbf{A}_{1,3} \\ \mathbf{A}_{2,0} & \mathbf{A}_{2,1} & \mathbf{A}_{2,2} & \mathbf{A}_{2,3} \\ \mathbf{A}_{3,0} & \mathbf{A}_{3,1} & \mathbf{A}_{3,2} & \mathbf{A}_{3,3} \end{pmatrix} \begin{pmatrix} \mathbf{x} \\ \mathbf{y} \\ \mathbf{z} \\ \mathbf{w} \end{pmatrix} = \begin{pmatrix} \mathbf{x}\mathbf{A}_{0,0} + \mathbf{y}\mathbf{A}_{0,1} + \mathbf{z}\mathbf{A}_{0,2} + \mathbf{w}\mathbf{A}_{0,3} \\ \mathbf{x}\mathbf{A}_{1,0} + \mathbf{y}\mathbf{A}_{1,1} + \mathbf{z}\mathbf{A}_{1,2} + \mathbf{w}\mathbf{A}_{1,3} \\ \mathbf{x}\mathbf{A}_{2,0} + \mathbf{y}\mathbf{A}_{2,1} + \mathbf{z}\mathbf{A}_{2,2} + \mathbf{w}\mathbf{A}_{2,3} \\ \mathbf{x}\mathbf{A}_{3,0} + \mathbf{y}\mathbf{A}_{3,1} + \mathbf{z}\mathbf{A}_{3,2} + \mathbf{w}\mathbf{A}_{3,3} \end{pmatrix}$$

Elements from input vector in every output row

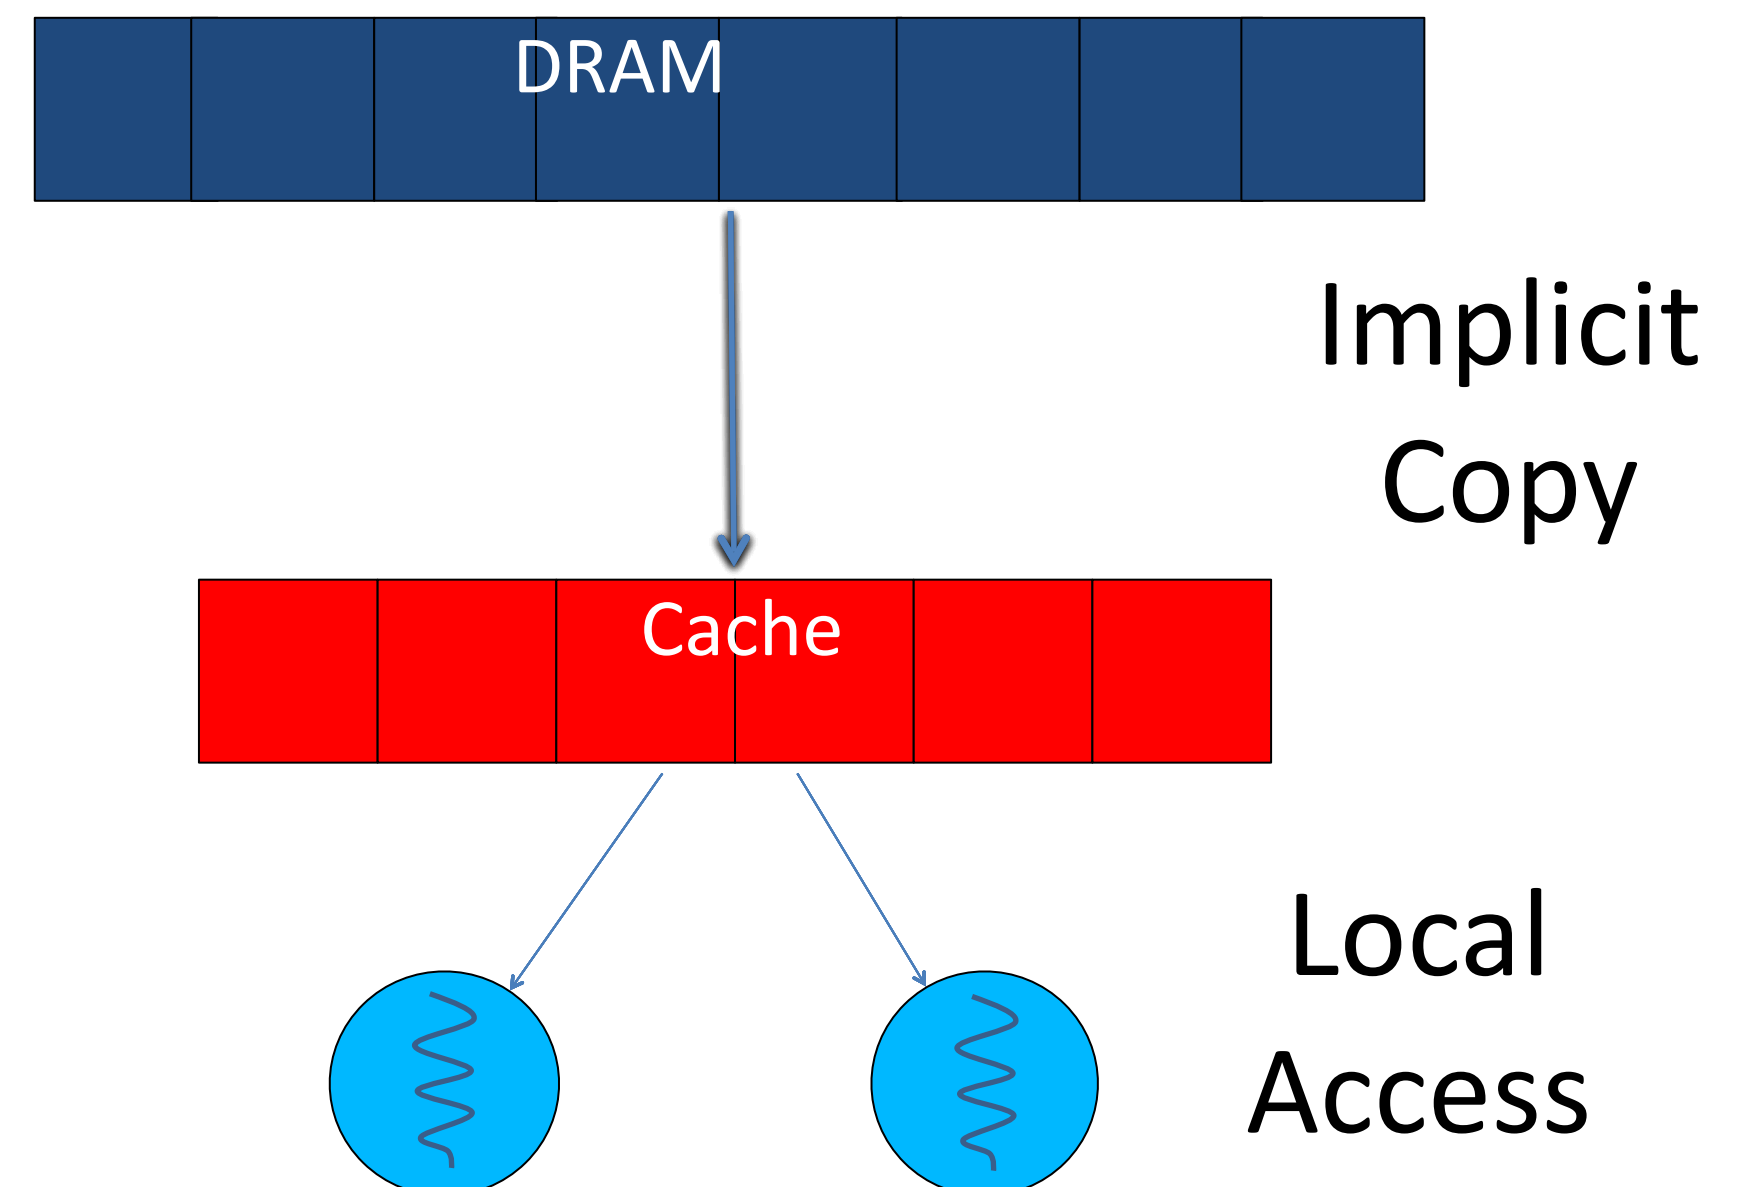
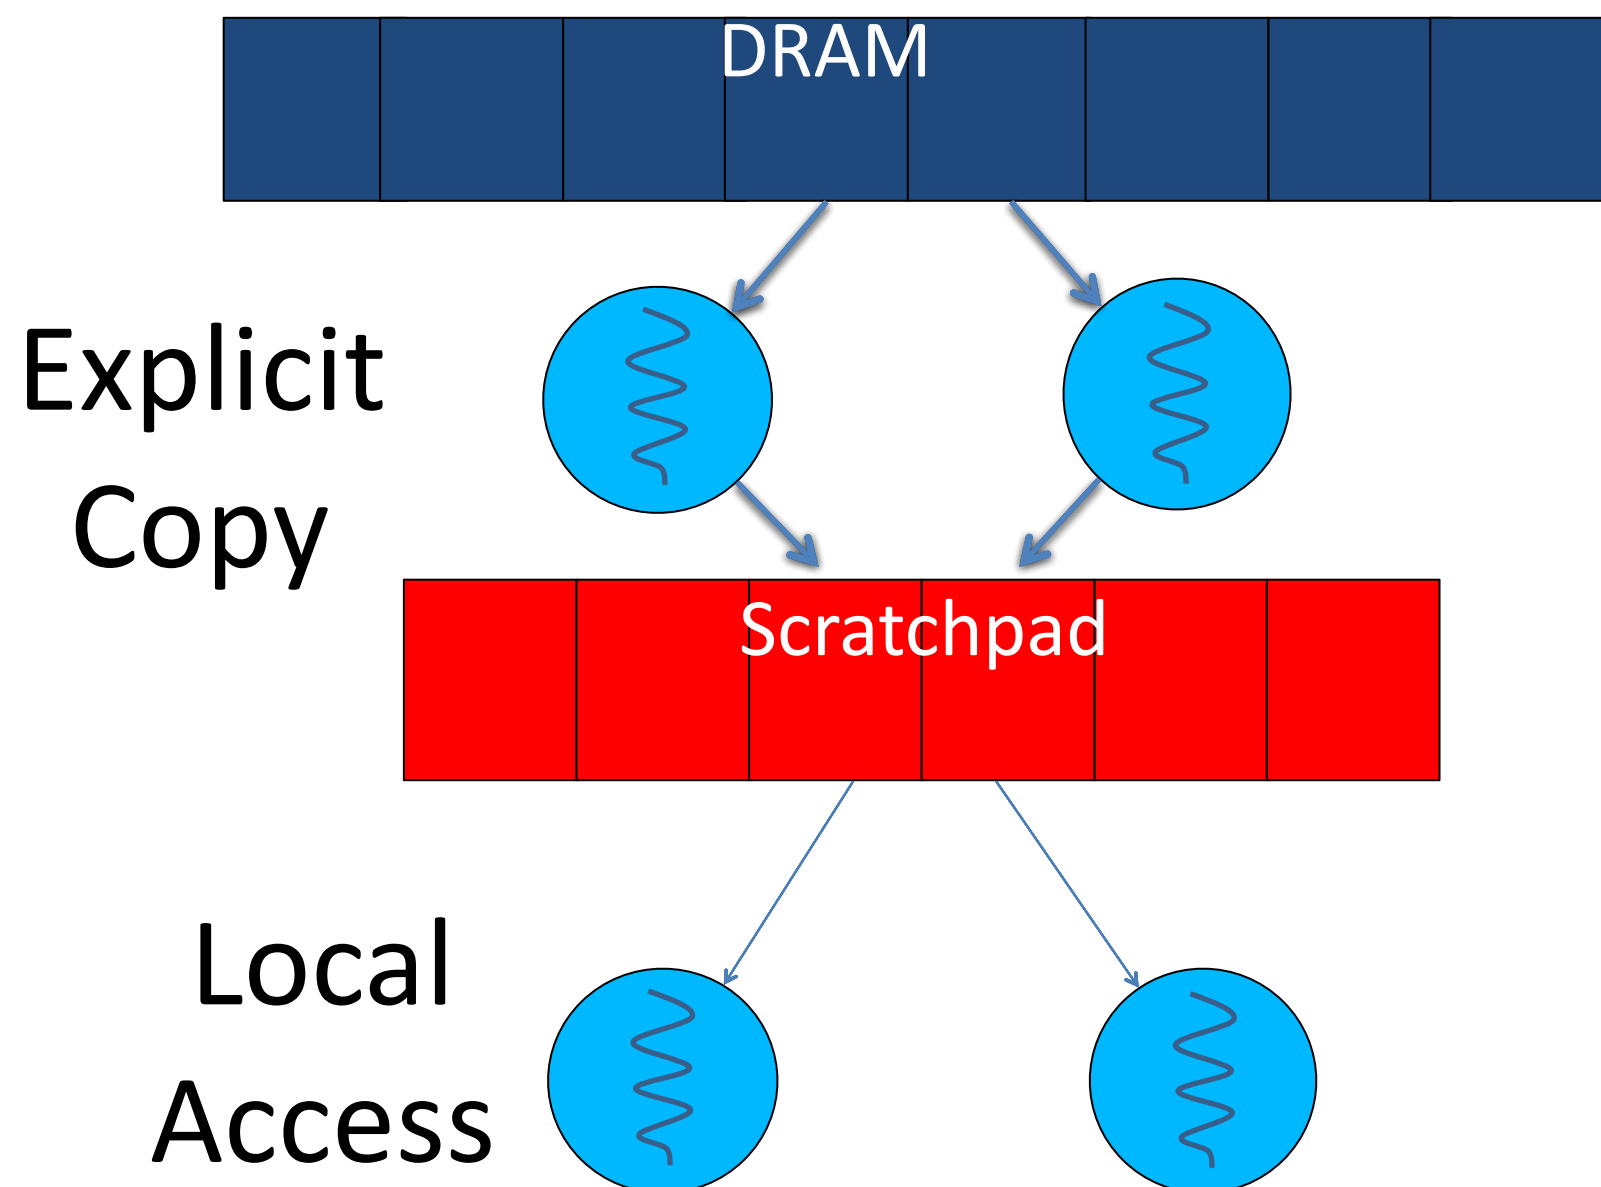

SpMV Kernel

```
__global__ void spmv(float **m, float *v, float *y) {  
    int row = threadIdx.x + blockIdx.x * blockDim.x;  
    int col = threadIdx.y + blockIdx.y * blockDim.y;  
  
    y[row] += m[row][col] * v[col];  
}
```

SpMV Kernel

```
__global__ void spmv(float **m, float *v, float *y) {  
    int row = threadIdx.x + blockIdx.x * blockDim.x;  
    int col = threadIdx.y + blockIdx.y * blockDim.y;  
  
    y[row] += m[row][col] * v[col];  
    atomicAdd(&y[row], m[row][col] * v[col]);  
}
```

1: (Input) Data Access Tiling



Data-Access Tiling

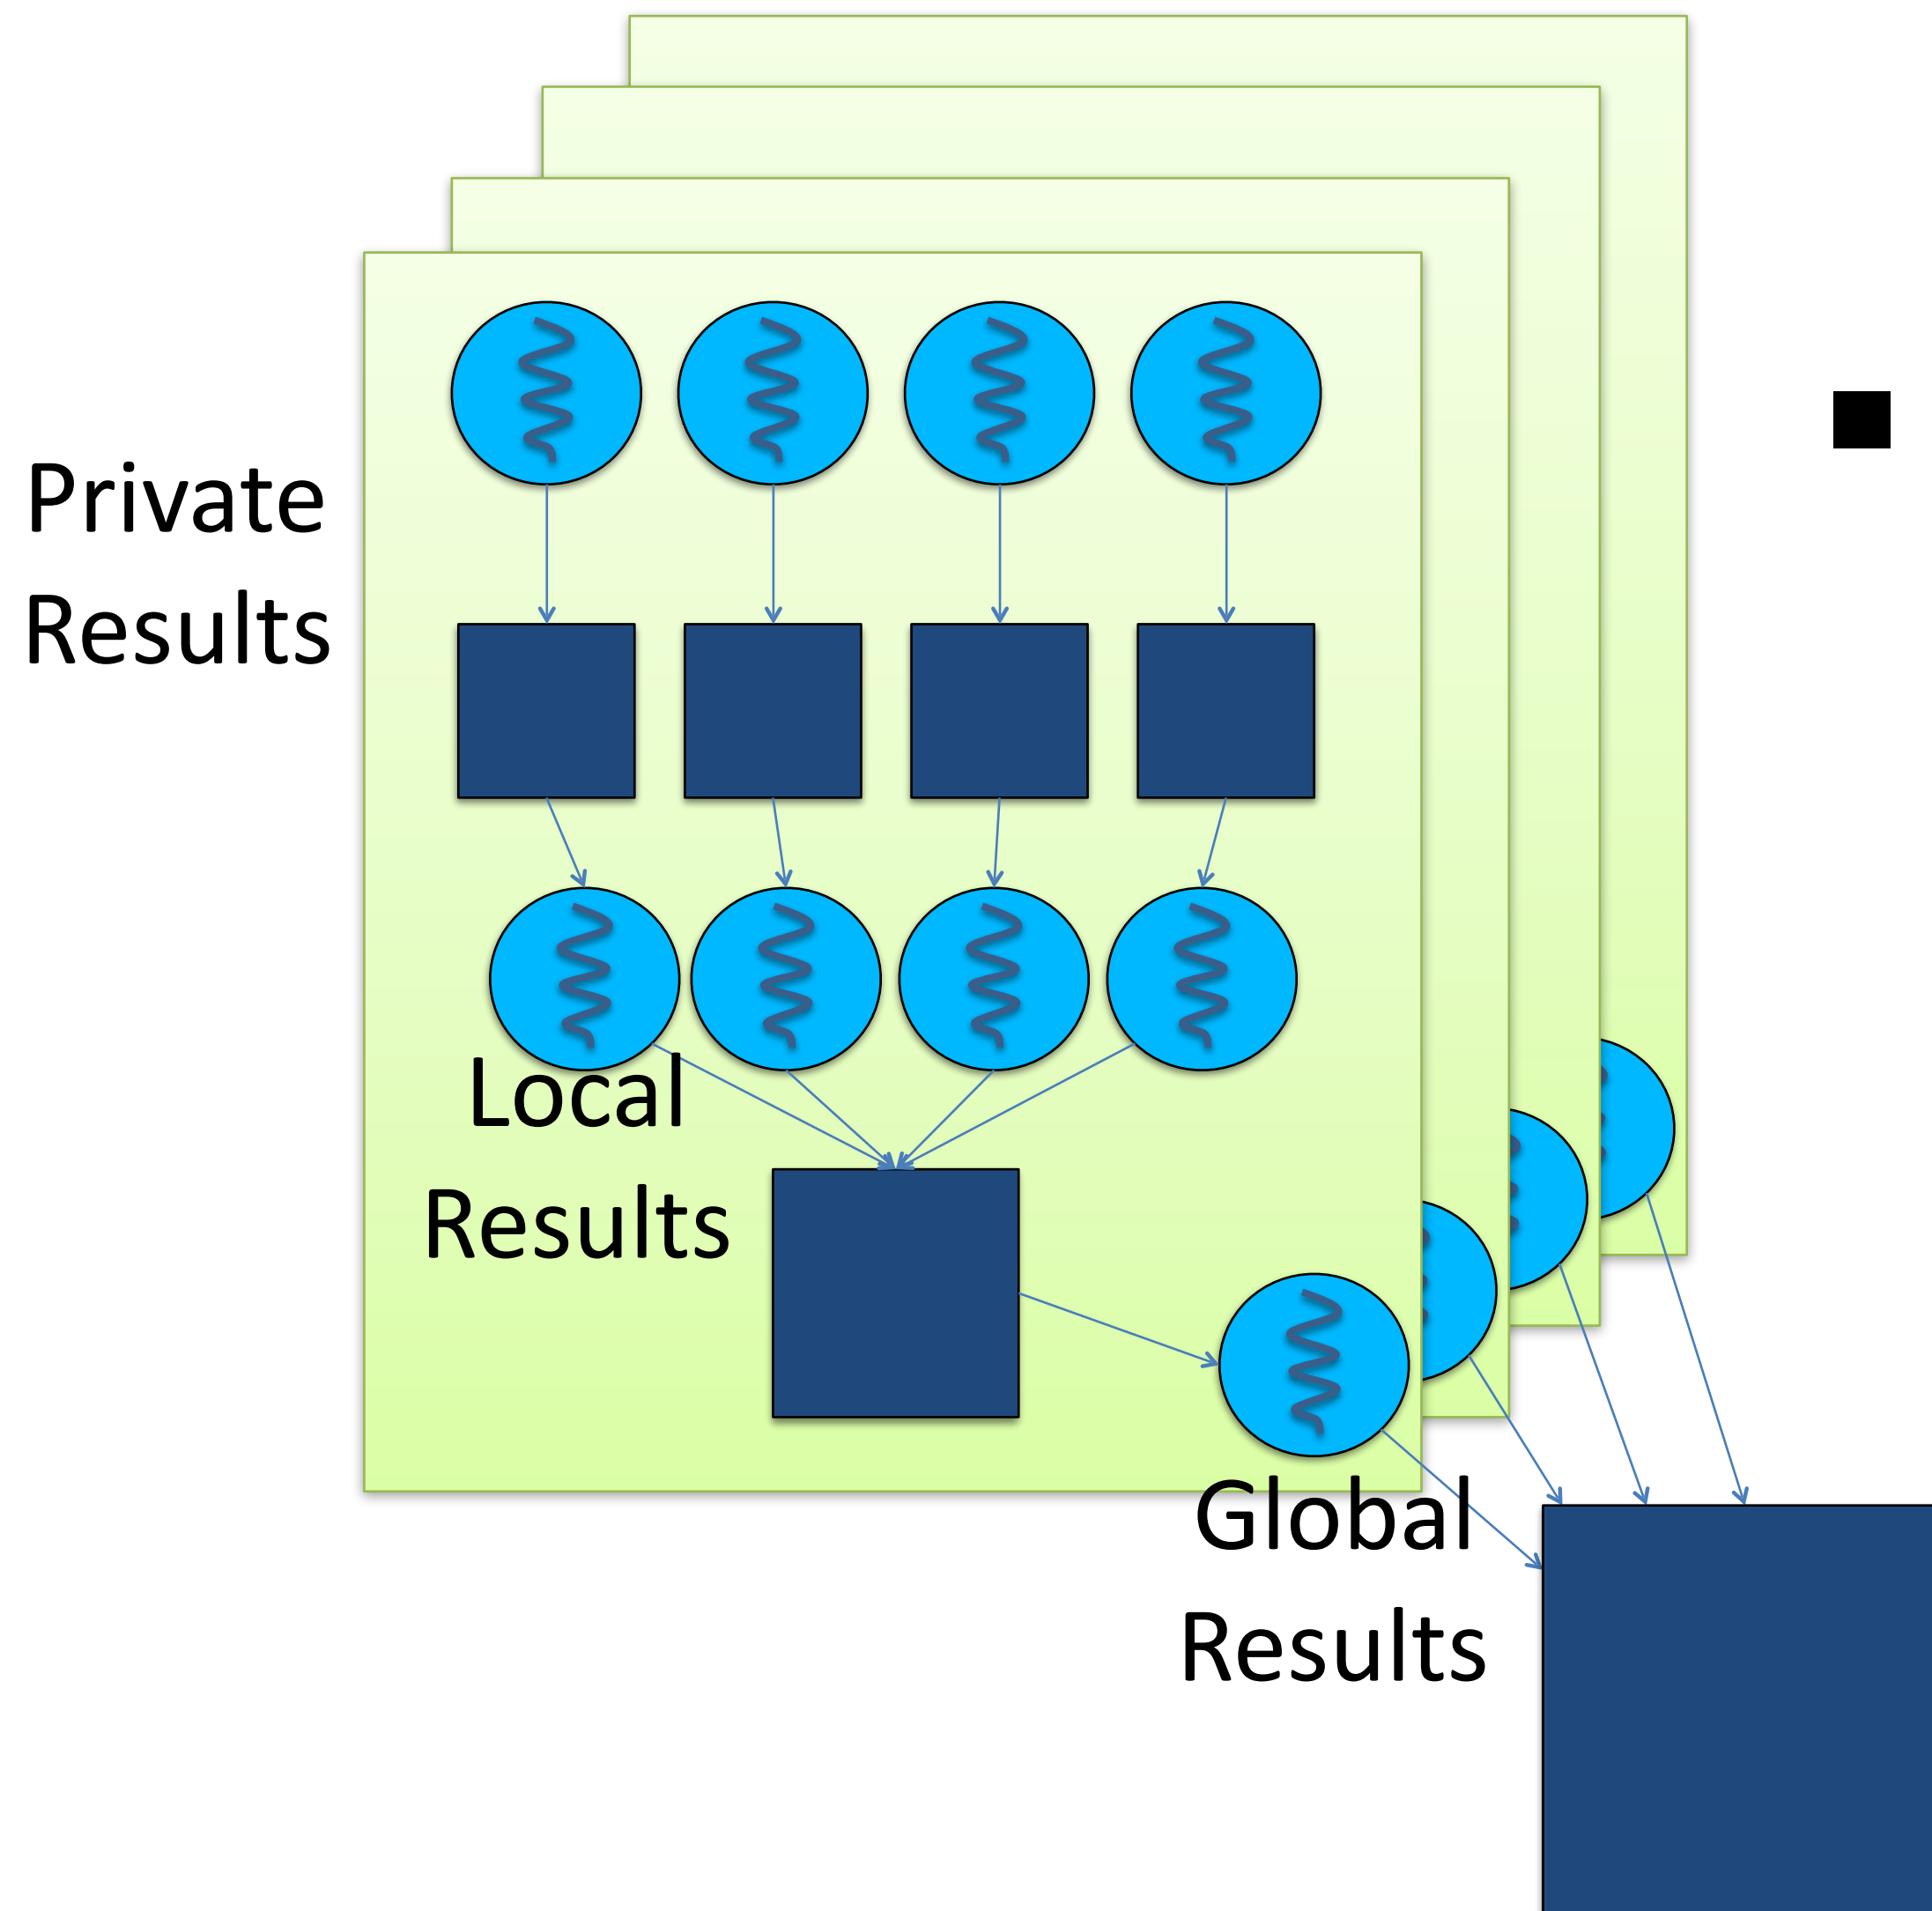
```
__global__ void spmv(float **m, float *v, float *y) {  
    int row = threadIdx.x + blockIdx.x * blockDim.x;  
    int col = threadIdx.y + blockIdx.y * blockDim.y;  
  
    __shared__ float vs[VECTOR_SIZE];  
    if(row == 0) {  
        vs[col] = v[col];  
    }  
    __syncthreads()  
  
    atomicAdd(&y[row], m[row][col] * vs[col]);  
}
```

1. (Input) Data Access Tiling

- **Pro: Better use of the memory system**
 - **Coalesced accesses**
 - **Data reuse**
- **Con: Reduced scheduling flexibility**
 - **Threads must synchronize**
 - **Larger shared memory use -> fewer blocks per SM**

2. (Output) Privatization

- Avoid contention by aggregating updates locally
- Requires storage resources to keep copies of data structures



Output Privatization

```
__global__ void spmv(float **m, float *v, float **yLocal) {  
    int row = threadIdx.x + blockIdx.x * blockDim.x;  
    int col  = threadIdx.y + blockIdx.y * blockDim.y;  
  
    yLocal[row][col] = m[row][col] * v[col];  
}
```

```
spmv<<<...>>>(m, v, yLocal);  
for(int row = 0; row < NUM_ROWS; ++row) {  
    y[row] = reduce("+", yLocal[row]);  
}
```

2. (Output) Privatization

- **Pro: Reduce write contention**
 - **Don't need atomics for every update**
- **Con: More memory usage**
 - **Need copy of data per thread**
- **Variant: One copy per block + smem atomics**
 - **smem atomics are faster on newer hardware**

Output Privatization (Variant)

```
__global__ void spmv(float **m, float *v, float *y) {  
    int row = threadIdx.x + blockIdx.x * blockDim.x;  
    int col  = threadIdx.y + blockIdx.y * blockDim.y;  
  
    __shared__ float ys[VECTOR_SIZE];  
    if(col == 0) {  
        ys[row] = 0;  
    }  
    __syncthreads()  
    atomicAdd(&ys[row], m[row][col] * v[col]); // Shared memory atomic  
    __syncthreads()  
    atomicAdd(&y[row], ys[row]); // Global memory atomic  
}
```

Storage Format Comparison

(DIA) Diagonal

(ELL) ELLPACK



(CSR) Compressed Row

(HYB) Hybrid

(COO) Coordinate

Structured

Unstructured



Slide credit: Nathan Bell

Coordinate Format (C00)

$$\begin{pmatrix} 3 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 2 & 4 & 1 \\ 1 & 0 & 0 & 1 \end{pmatrix}$$

Coordinate Format (C00)

$$\begin{pmatrix} 3 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 2 & 4 & 1 \\ 1 & 0 & 0 & 1 \end{pmatrix}$$

Non-zero values = $\{ 3, 1, 2, 4, 1, 1, 1 \}$

Row indices = $\{ 0, 0, 2, 2, 2, 3, 3 \}$

Column indices = $\{ 0, 2, 1, 2, 3, 0, 3 \}$

Coordinate Format (C00)

$$\begin{pmatrix} 3 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 2 & 4 & 1 \\ 1 & 0 & 0 & 1 \end{pmatrix}$$


Non-zero values = {3, 1, 2, 4, 1, 1, 1}

Row indices = {0, 0, 2, 2, 2, 3, 3}

Column indices = {0, 2, 1, 2, 3, 0, 3}


Running Example: SpMV

$$Ax = v$$




Running Example: SpMV

$$Ax = v$$




Running Example: SpMV

$$Ax = v$$




Running Example: SpMV

$$Ax = v$$



Running Example: SpMV

$$Ax = v$$




3. “Scatter to Gather” Transformation


- **Write conflicts have to be serialized (atomics)**
- **Turn overlapping writes into overlapping reads**
- **Hardware can handle overlapping reads more efficiently**

3. "Scatter to Gather" Transformation


$$Ax = v$$




3. "Scatter to Gather" Transformation



3. "Scatter to Gather" Transformation



3. "Scatter to Gather" Transformation



3. “Scatter to Gather” Transformation


- **Pro: Reduce write contention**
 - **Don't need atomics for every update**
- **Con: Harder to program efficiently**
 - **Harder to get memory coalescing**

4. Binning

- **Usually easy to find scatter address from input**
 - **E.g., Matrix row number \rightarrow output location**
- **Often difficult to find gather addresses efficiently**
 - **E.g., For sparse matrices, read a lot of unnecessary data**


3. "Scatter to Gather" Transformation

$$Ax = v$$




3. "Scatter to Gather" Transformation

$$Ax = v$$




3. "Scatter to Gather" Transformation

x



$$Ax = v$$


v



Row

Col


Data




A

3. "Scatter to Gather" Transformation






$$Ax = v$$



4. Binning




4. Binning



A

4. Binning



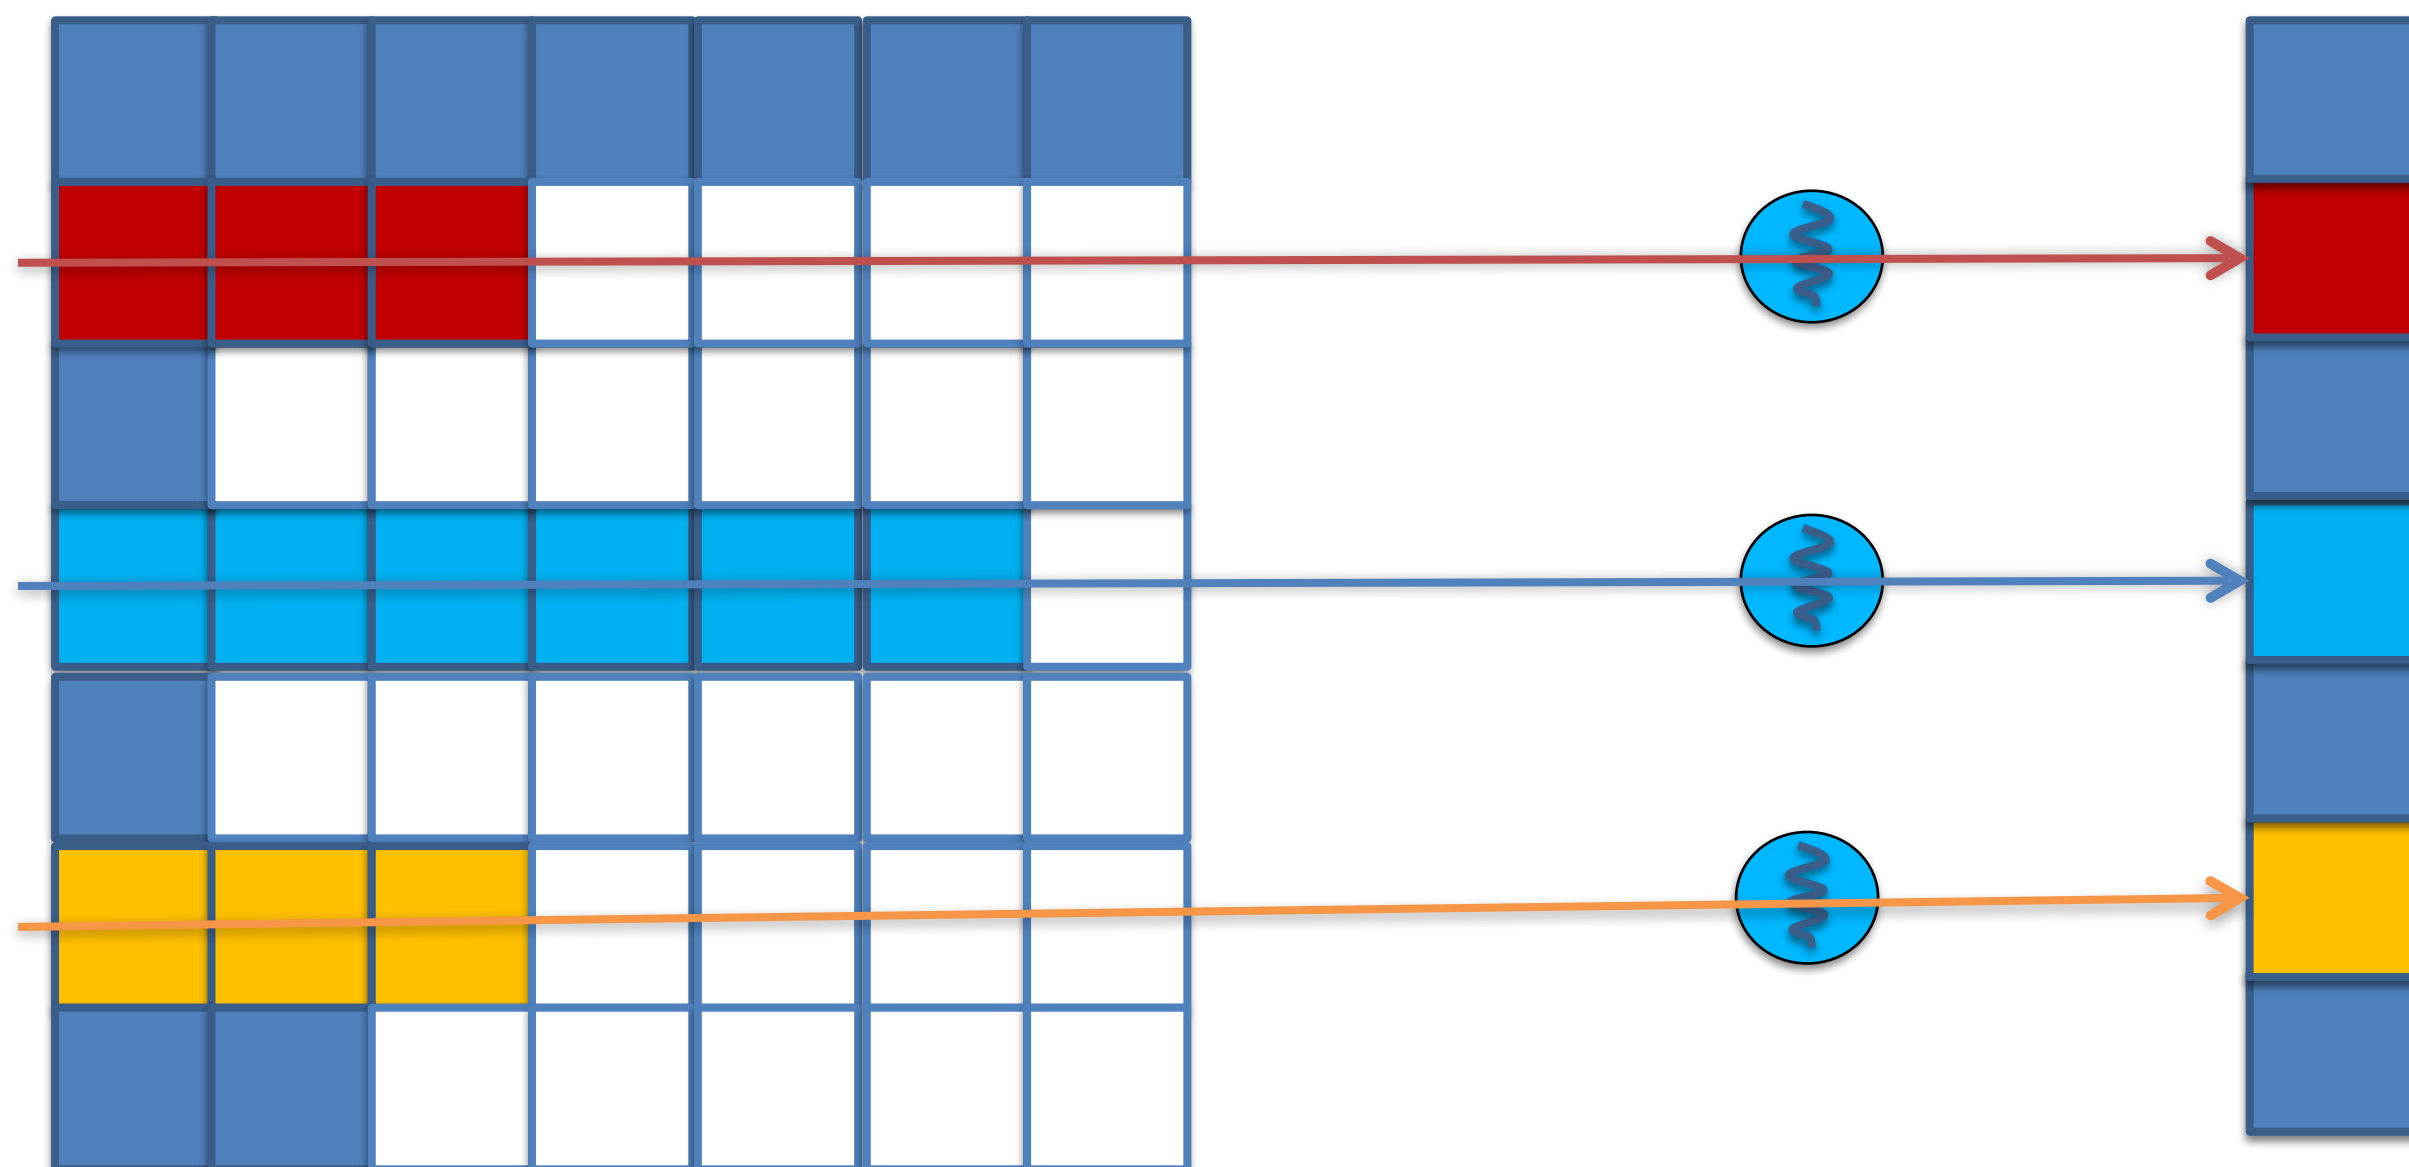
A

4. Binning

- **Pro: Organize data for easy access**
 - **Without accessing unnecessary data**
- **Con: Large memory requirement**
 - **Lots of wasted space (How can we fix this?)**
- **Con: Binning takes time**


5. Compaction

$$\begin{pmatrix} 3 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 2 & 4 & 1 \\ 1 & 0 & 0 & 1 \end{pmatrix}$$



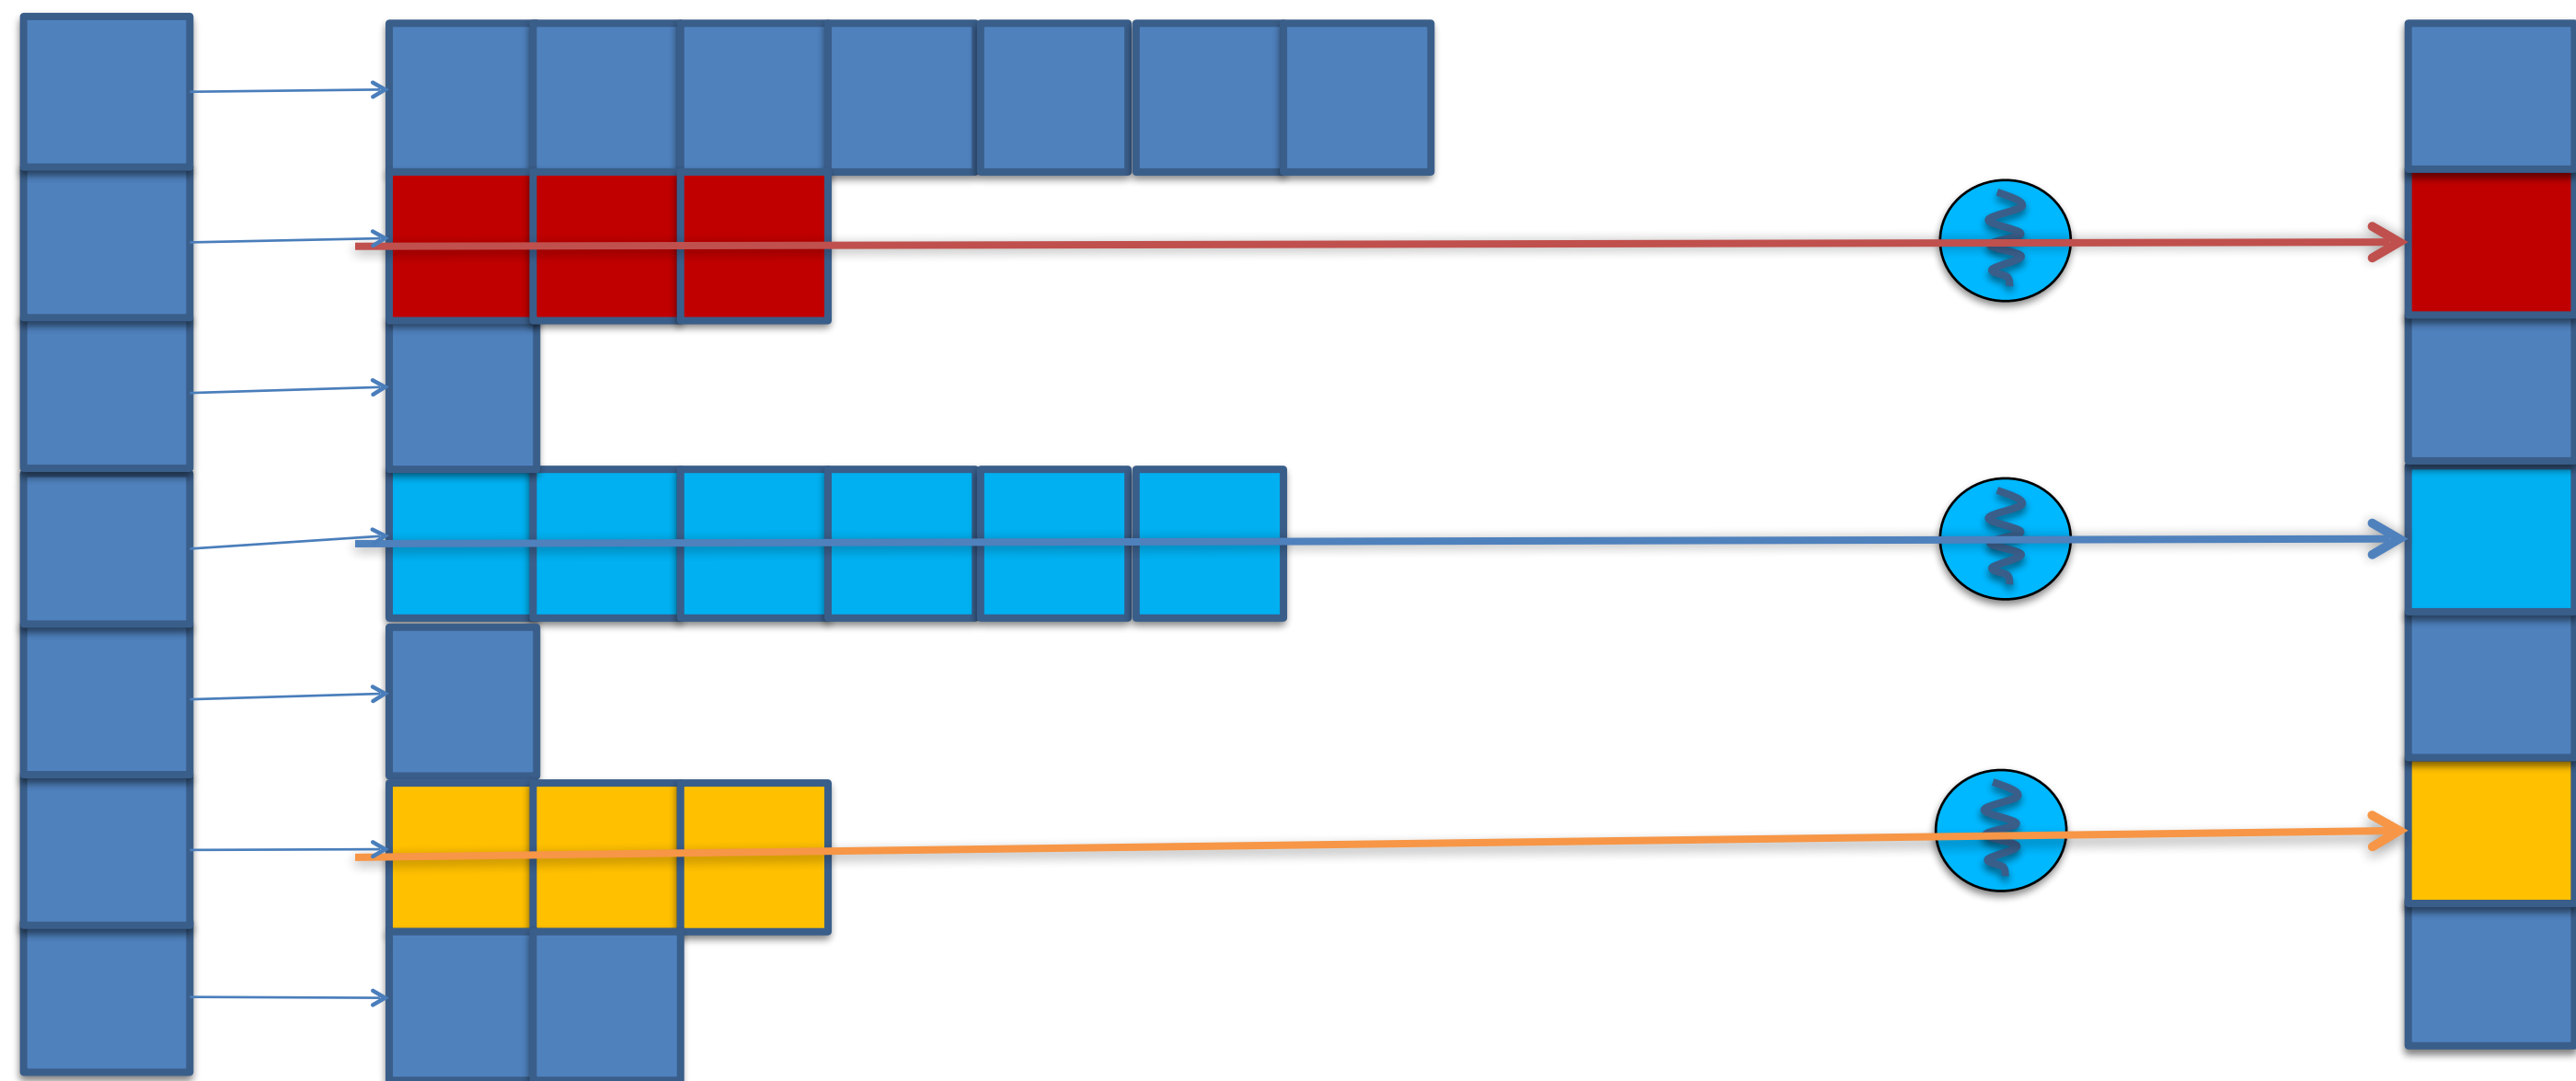
5. Compaction

$$\begin{pmatrix} 3 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 2 & 4 & 1 \\ 1 & 0 & 0 & 1 \end{pmatrix}$$



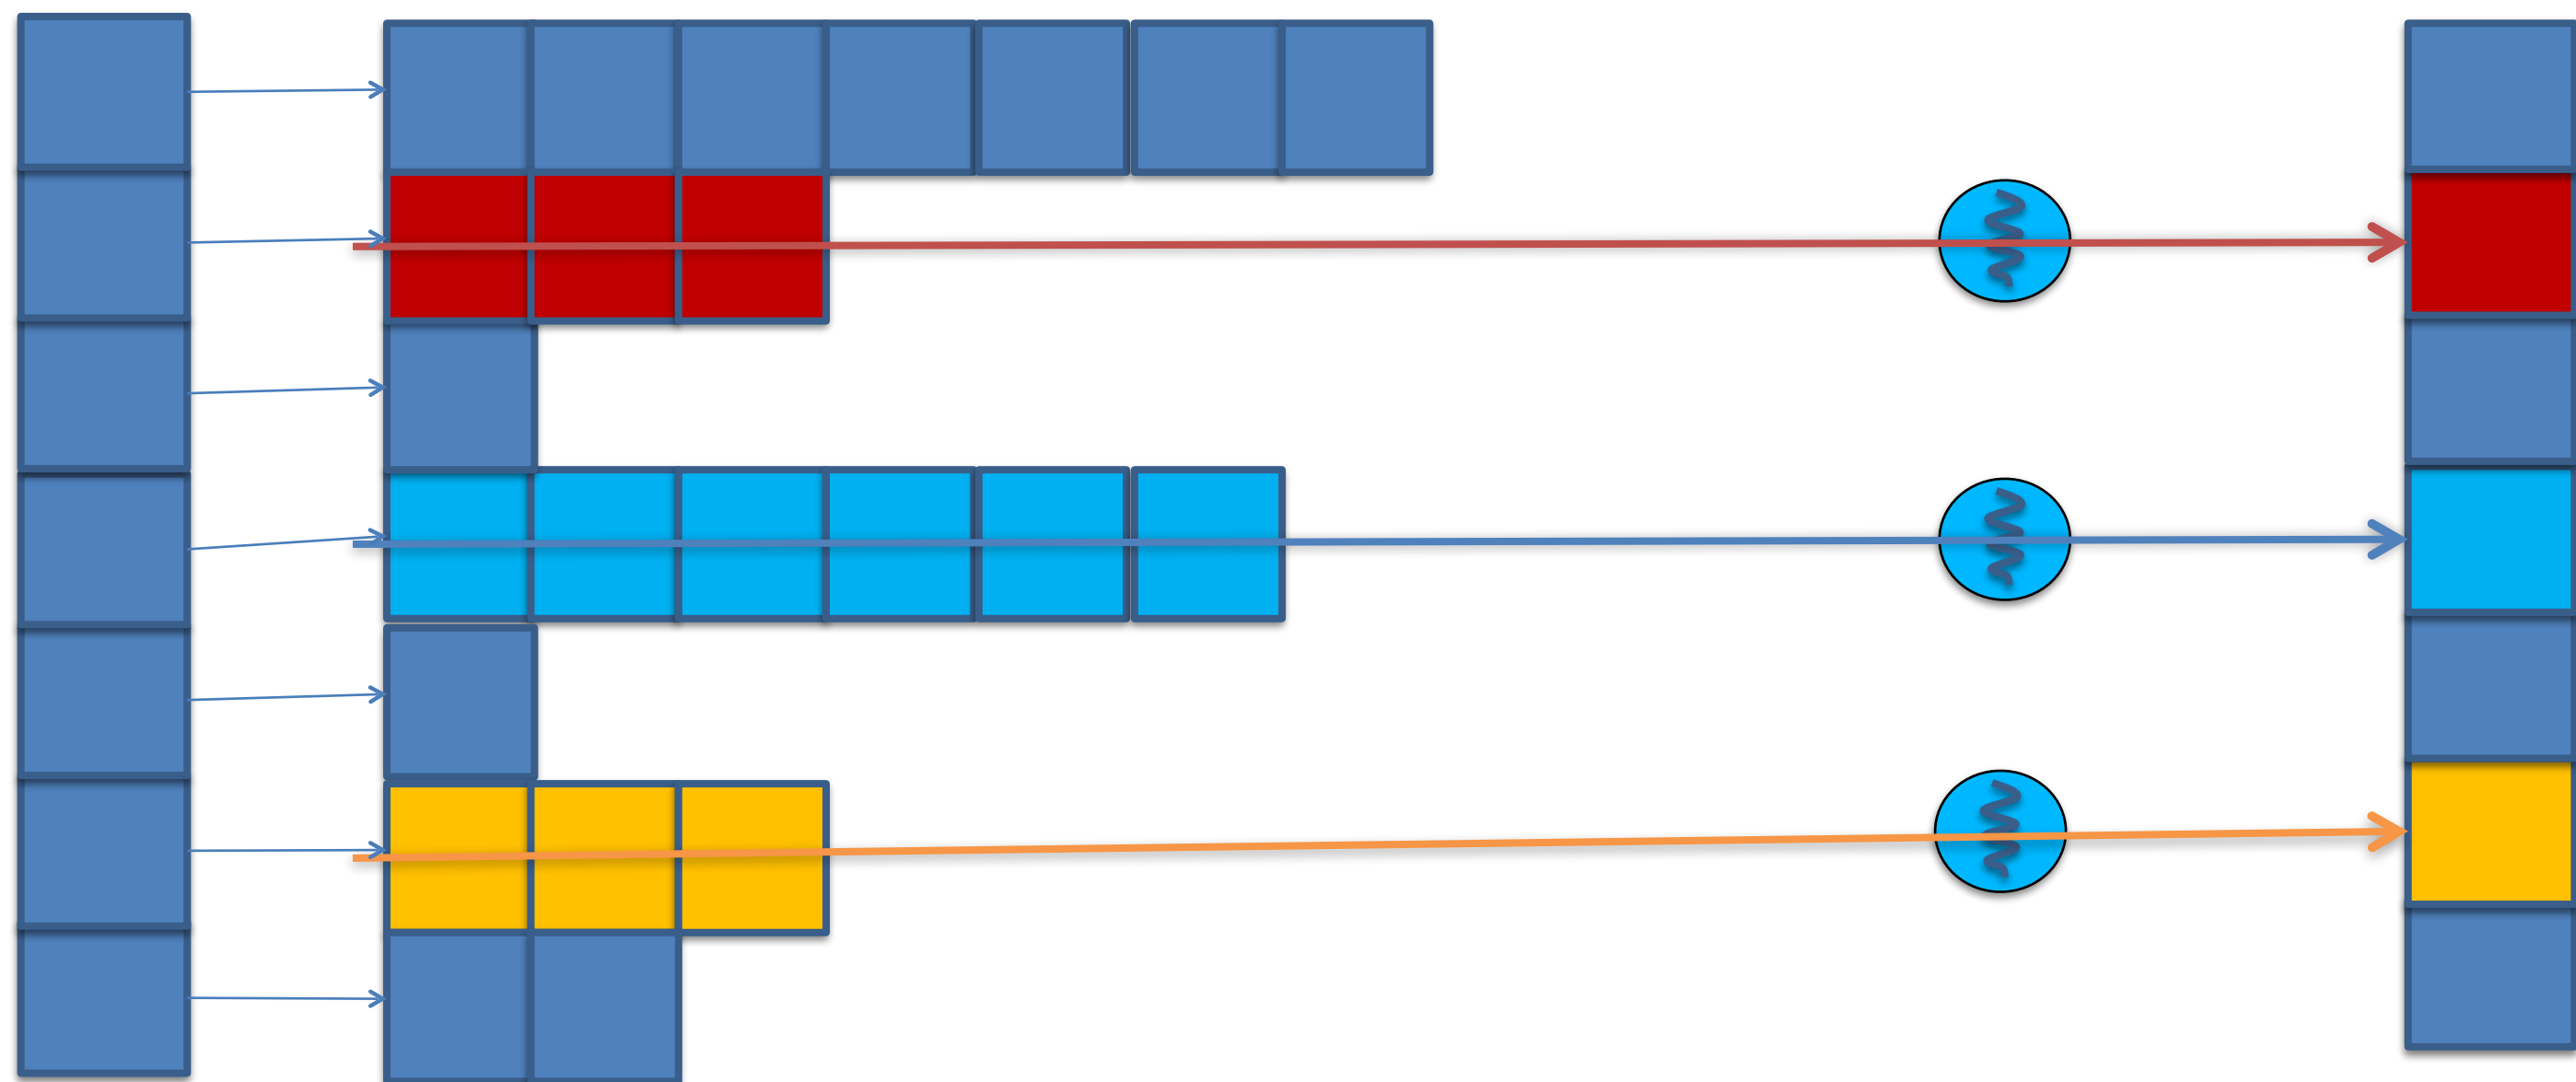
5. Compaction

$$\begin{pmatrix} 3 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 2 & 4 & 1 \\ 1 & 0 & 0 & 1 \end{pmatrix}$$



5. Compaction


3	0	1	0
0	0	0	0
0	2	4	1
1	0	0	1




6. Regularization (Load Balancing)

- **Some threads have more work than others**
 - **Poor hardware utilization if we're waiting for a few threads to finish**
 - **Can't release resources until all threads in block have finished**
- **Want to balance the work between threads**

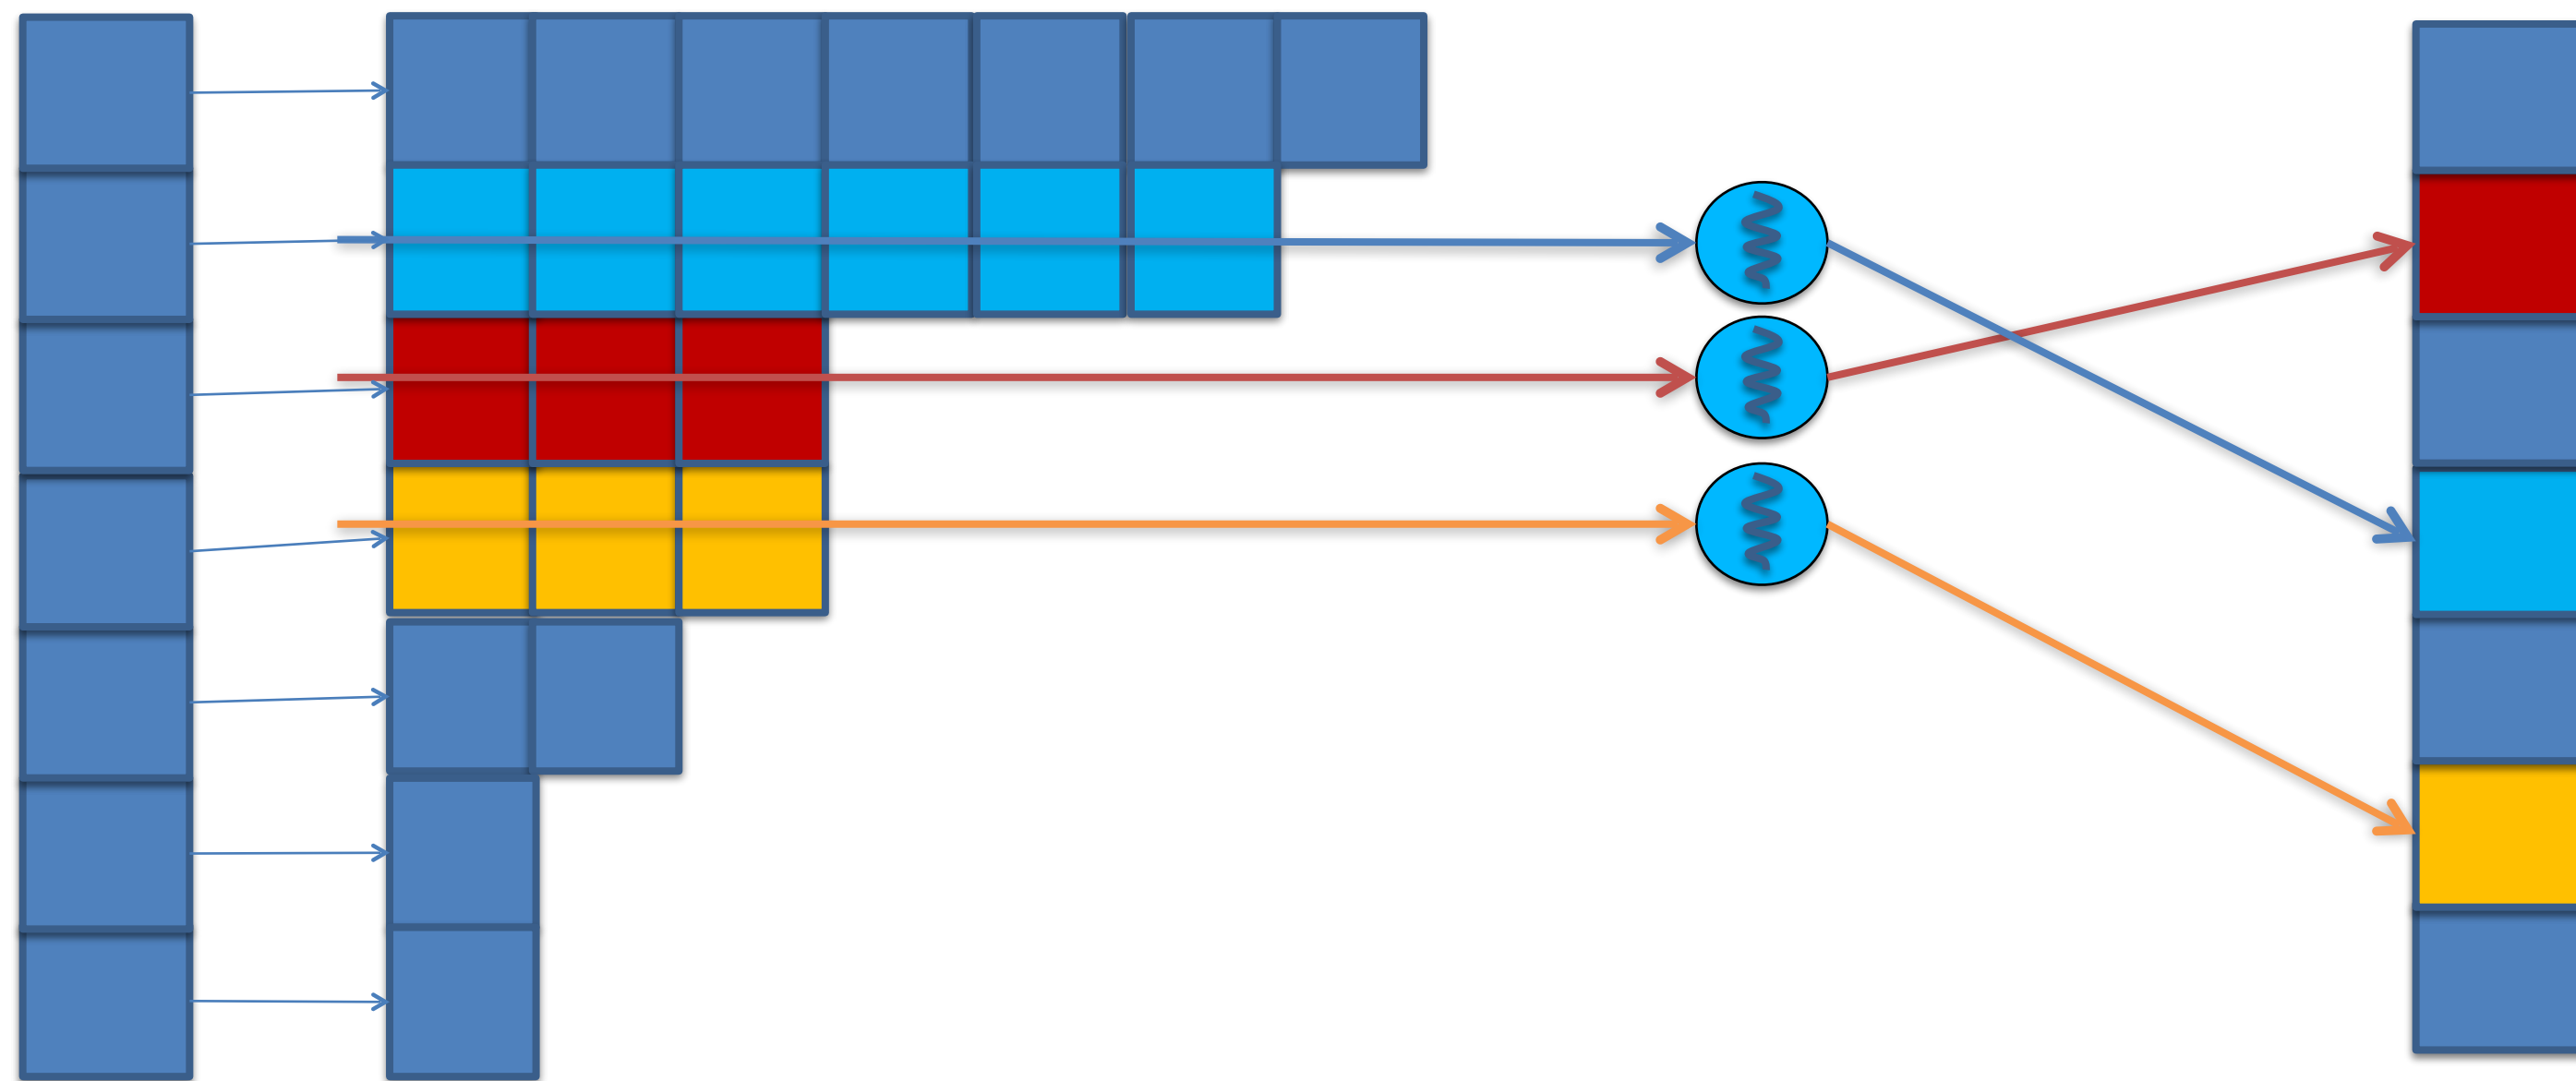
6. Regularization (Load Balancing)



6. Regularization (Load Balancing)




6. Regularization (Load Balancing)




6. Regularization (Load Balancing)

- **Pro: Better hardware utilization**
- **Con: Some applications difficult to load balance**
- **Con: Load balancing operations take finite time**

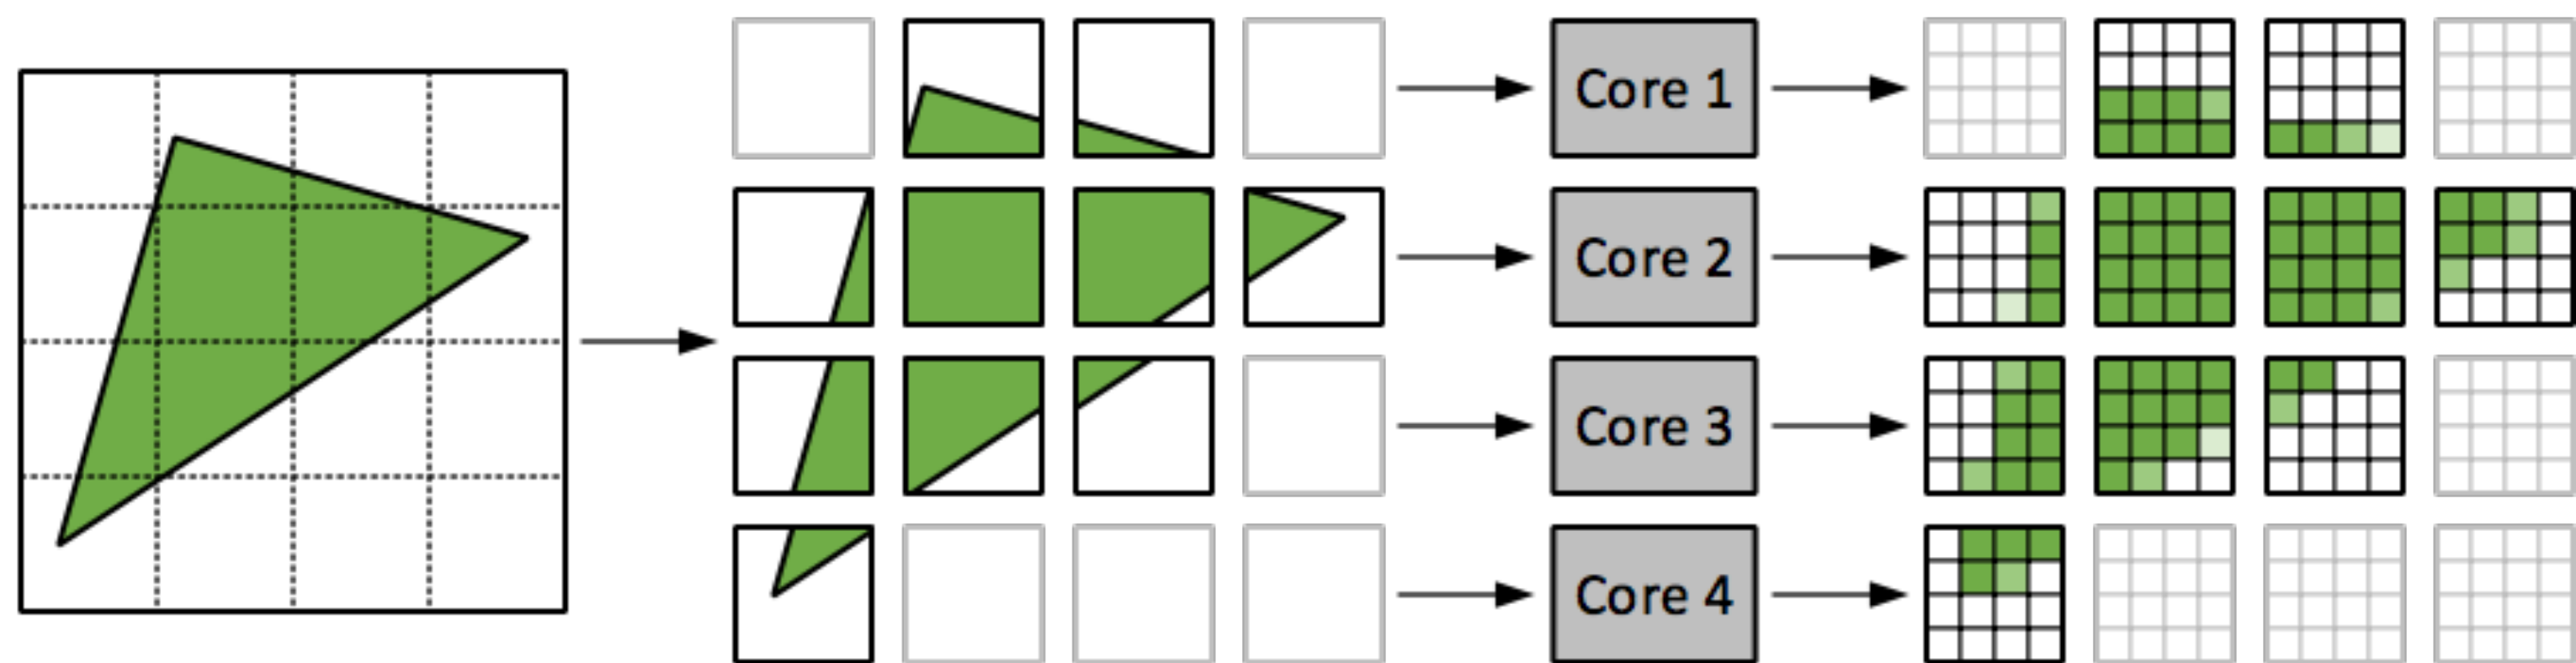
Piko: Spatial Tiling for Parallelism



Piko: Spatial Tiling for Parallelism



Piko: Spatial Tiling for Parallelism





7. Data Layout Transformation

■ “Array of Structs” vs. “Struct of Arrays”

```
struct Data {  
    float a;  
    float b;  
    float c;  
};
```

```
struct Data {  
    float a[];  
    float b[];  
    float c[];  
};
```



7. Data Layout Transformation



■ “Array of Structs” vs. “Struct of Arrays”

```
struct Data {  
    float a;  
    float b;  
    float c;  
};
```

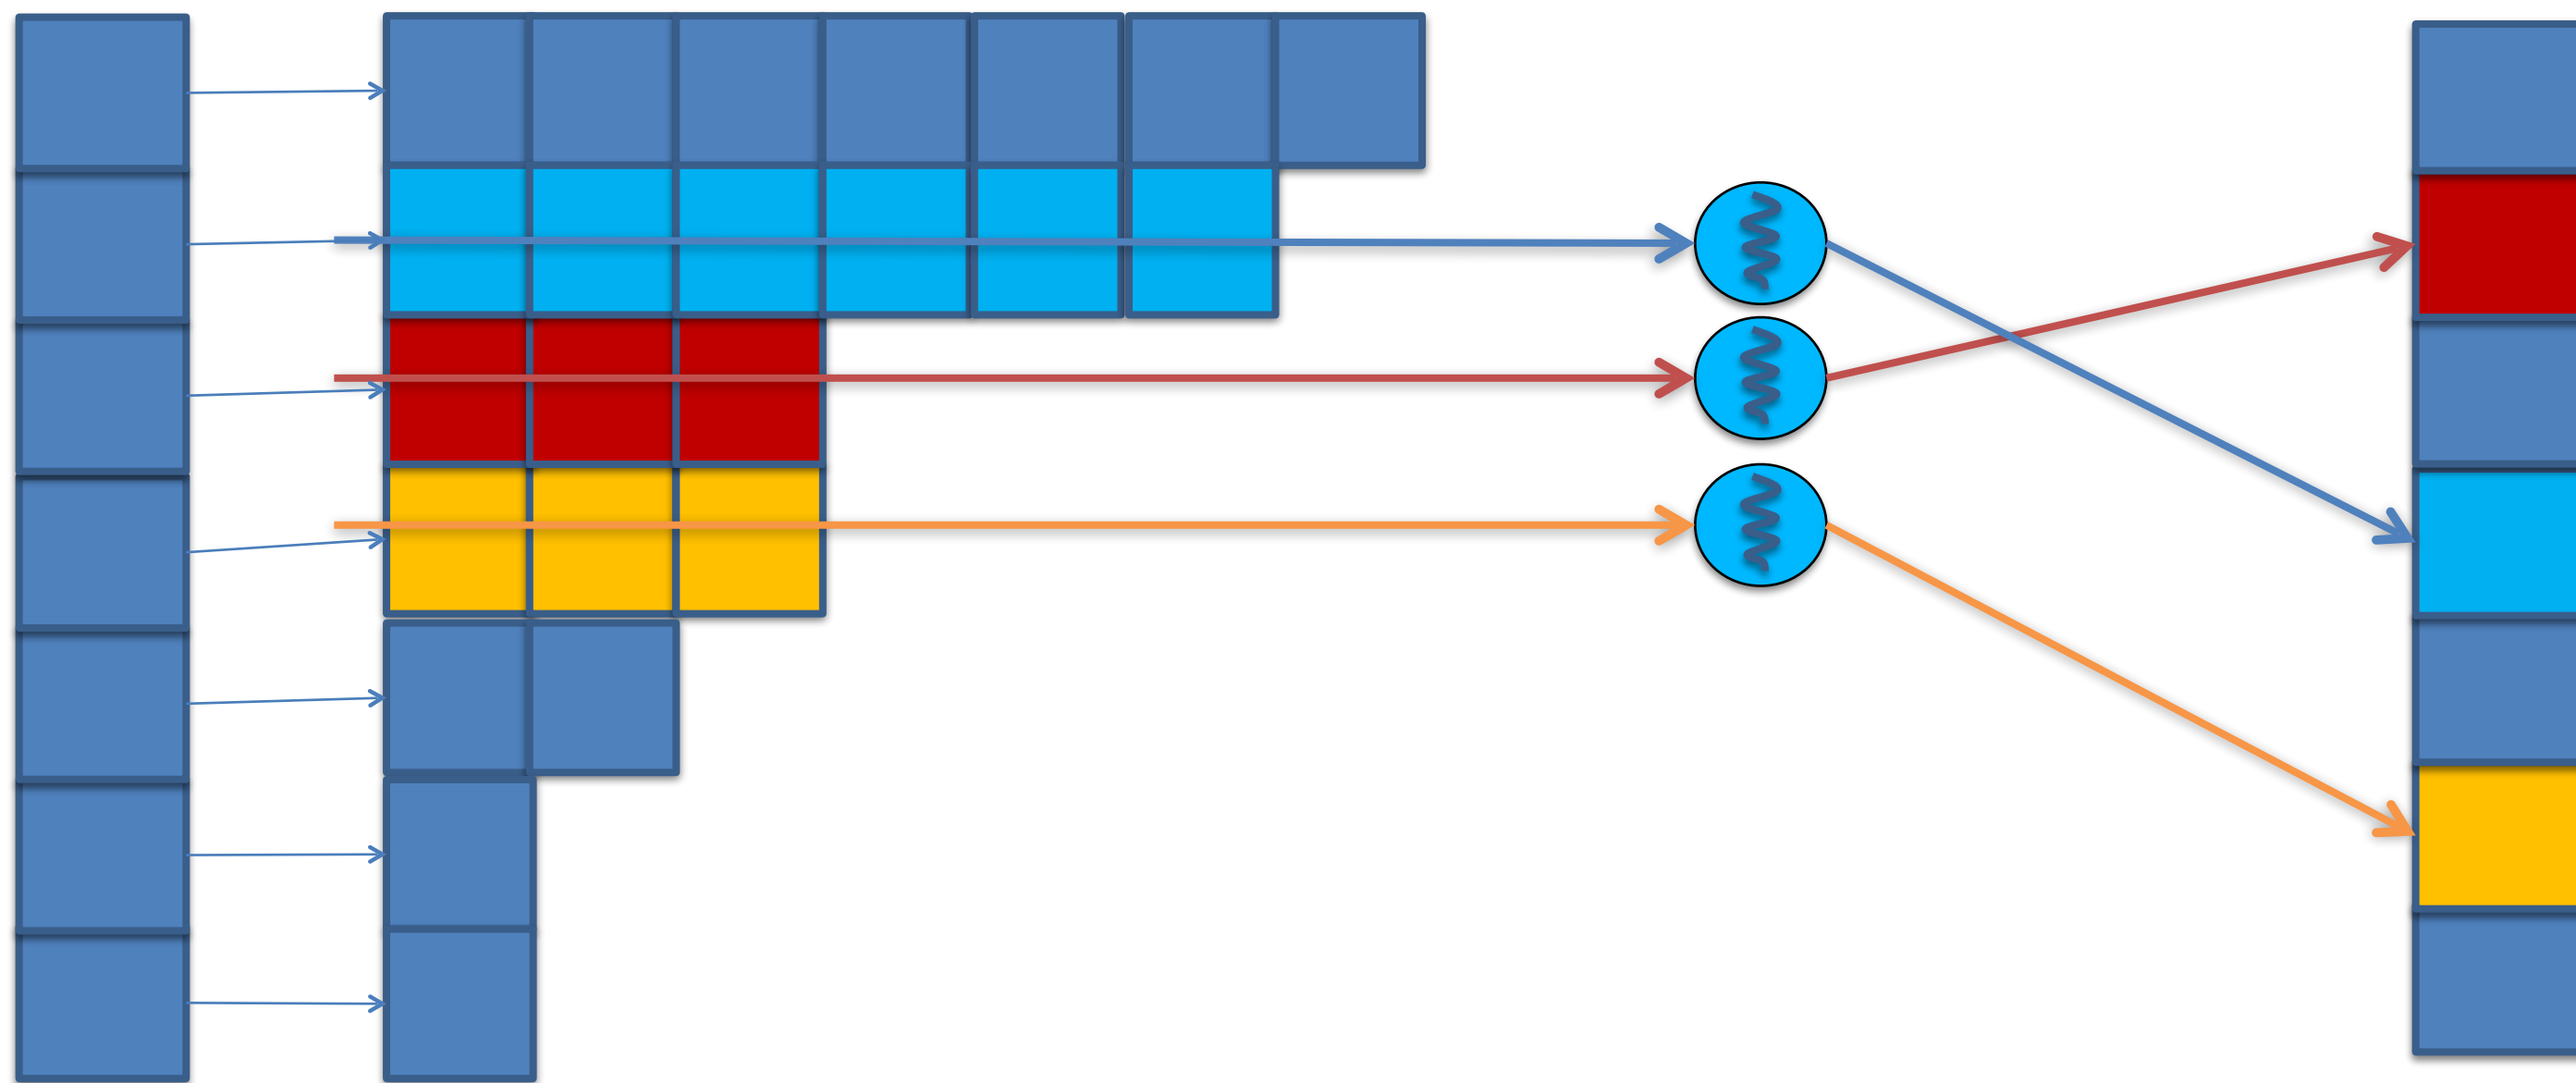
Better for
CPU

```
struct Data {  
    float a[];  
    float b[];  
    float c[];  
};
```


Better for
GPU



7. Data Layout Transformation



7. Data Layout Transformation




7. Data Layout Transformation

- **Pro: Better memory access patterns for GPU**
- **Con: Usually requires reorganizing data (takes time)**
 - **Beautiful piece of work from NVIDIA called “Trove” that does a matrix transpose on the fly**
 - **Many pieces of work that discuss $AOS \leftrightarrow SOA$**

8. Granularity Coarsening

- **Parallel execution often requires redundant work**
- **Let each thread process >1 element to reduce redundancy**

8. Granularity Coarsening




8. Granularity Coarsening



- **Pro: Reduces redundant computation**
- **Con: Reduces parallelism**
 - **And, thus, latency hiding potential**

9. Kernel Fusion

9. Kernel Fusion



9. Kernel Fusion



9. Kernel Fusion

- **Pro: Removes unnecessary reads/writes**
- **Con: Might lead to load imbalance**
 - **(when different threads generate different amounts of intermediate data/work)**

Optimization Summary

- **(Input) Data Access Tiling**
- **(Output) Privatization**
- **“Scatter to Gather” Transformation**
- **Binning**
- **Compaction**
- **Regularization (Load Balancing)**
- **Data Layout Transformation**
- **Granularity Coarsening**
- **Kernel Fusion**